

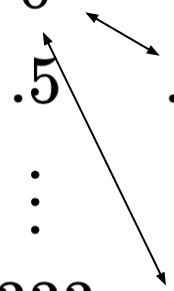
Overview

- What/Why is a Convolution?
- CNN-specific hyperparameters
- Basic CNN history/set-up

Why are images special?

- Images are deceptively hard
- Images are big
- Geometry matters!
 - Pixels near each other interact in different ways to create features than pixels far away
 - This is free data that we lose if we simply consider an image as a data vector

Different
relationships

$$\begin{bmatrix} 0 & 0 & 0 & \dots & 0 \\ .5 & .75 & 1 & \dots & .25 \\ \vdots & \vdots & \vdots & & \vdots \\ .333 & 0 & 1 & \dots & 0 \end{bmatrix}$$


The Convolution

- Fancy **linear** operation useful for spatial data

The Convolution

- Fancy **linear** operation useful for spatial data

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*

The Convolution

- Fancy **linear** operation useful for spatial data

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

*

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

=

$$\begin{bmatrix} & & & \\ & & & \\ & & & \\ & & & \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

$$(1 \times 1) + (.5 \times 0) + (0 \times 0) + (.25 \times 2) \\ = 1.5$$

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*

=

$$\begin{bmatrix} 1.5 & \dots & \dots \\ \vdots & \ddots & \vdots \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

$$(.5 \times 1) + (1 \times 0) + (.25 \times 0) + (.5 \times 2) = 1.5$$

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*

$$= \begin{bmatrix} 1.5 & \boxed{1.5} \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

$$(1 \times 1) + (0 \times 0) + (.5 \times 0) + (1 \times 2) = 3$$

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*

$$= \begin{bmatrix} 1.5 & 1.5 & \rightarrow 3 \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

$$(0 \times 1) + (.25 \times 0) + (1 \times 0) + (.25 \times 2) = .5$$

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*



$$\begin{bmatrix} 1.5 & 1.5 & 3 \\ 0.5 & & \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

*

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

=

$$\begin{bmatrix} 1.5 & 1.5 & 3 \\ 0.5 & ? & ? \\ ? & ? & ? \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

*

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

=

$$\begin{bmatrix} 1.5 & 1.5 & 3 \\ 0.5 & .25 & 2.5 \\ 1 & 2.25 & 2 \end{bmatrix}$$

The Convolution


- Fancy **linear** operation useful for spatial data
- Element-wise product

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{Result}$$

Unknown Parameters

“2x2 Filter”



Why Convolution?

- Only four parameters!
 - If input is dimension 16 and output is dimension 9, how many for FC?

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{Output Image}$$

Unknown Parameters

“2x2 Filter”

Why Convolution?


- Only four parameters!
- Translational Equivariance
 - If I shift my image, I shift the output!

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{Image}$$

Unknown Parameters

“2x2 Filter”



Why Convolution?

- Only four parameters!
- Translational Equivariance
- Weight Sharing (detect same feature translated to different parts of the image)

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{[Output Box]}$$

Unknown Parameters

“2x2 Filter”

Why Convolution?

Intuition: Edge
Detection

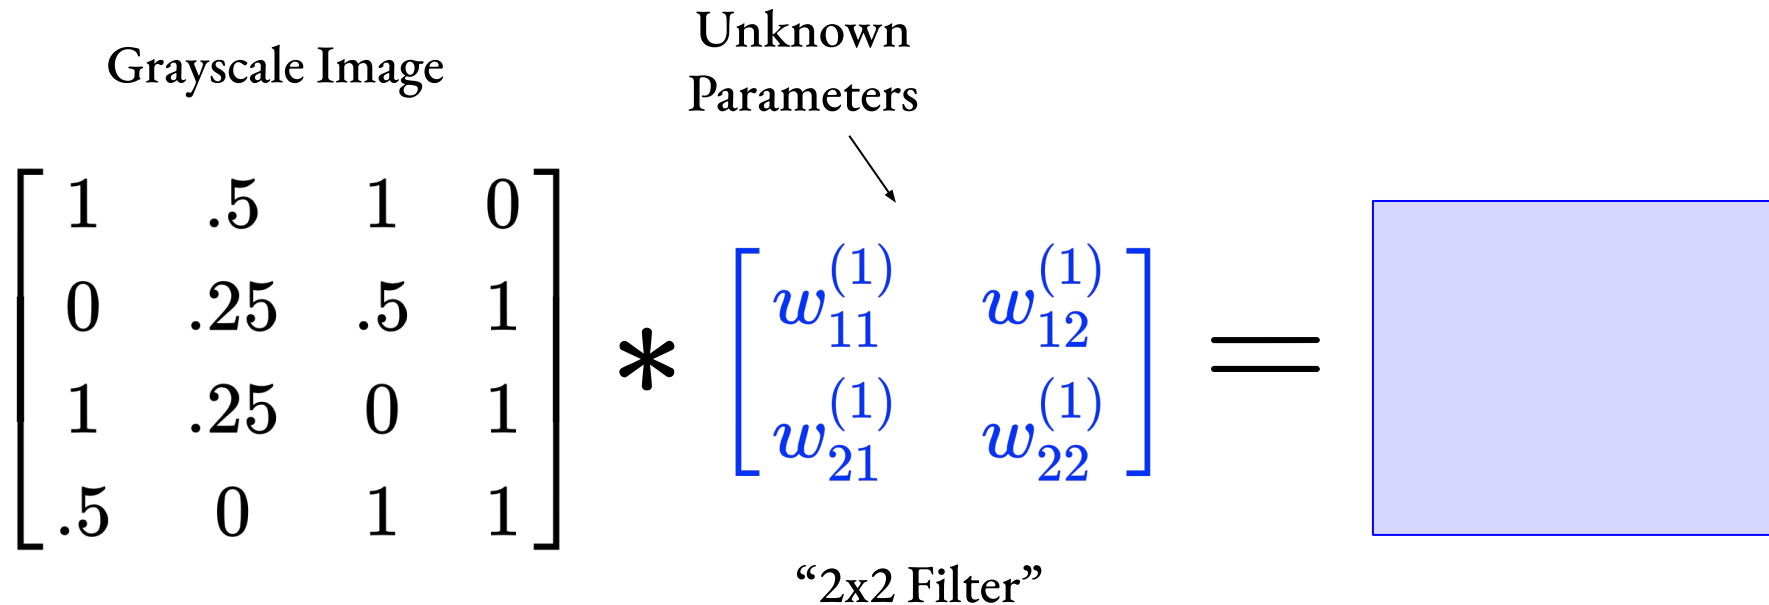
- Only four parameters!
- Translational Equivariance
- Weight Sharing (detect same feature translated to different parts of the image)

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{[Output Image]}$$

Unknown Parameters

“2x2 Filter”



The Convolution

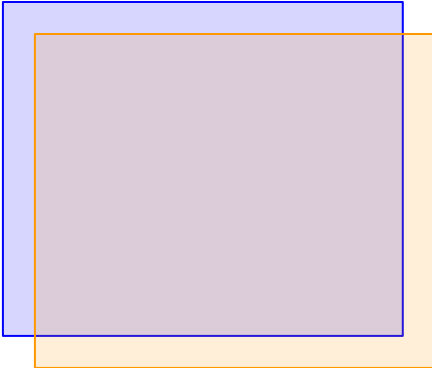
- In a Conv. layer we apply many filter to get many features

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(2)} & w_{12}^{(2)} \\ w_{21}^{(2)} & w_{22}^{(2)} \end{bmatrix} =$$

More Parameters

“2x2 Filter”



The diagram illustrates a 2D convolution operation. On the left, a 4x4 grayscale image matrix is shown. This matrix is convolved (indicated by the asterisk) with a 2x2 filter matrix, labeled "2x2 Filter" and "More Parameters". The filter matrix contains elements $w_{11}^{(2)}$, $w_{12}^{(2)}$, $w_{21}^{(2)}$, and $w_{22}^{(2)}$. An arrow points from the text "More Parameters" to the filter matrix. The result of the convolution is a 3x3 feature map, represented by a blue-outlined square with a purple-filled center and an orange-outlined border.

The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”

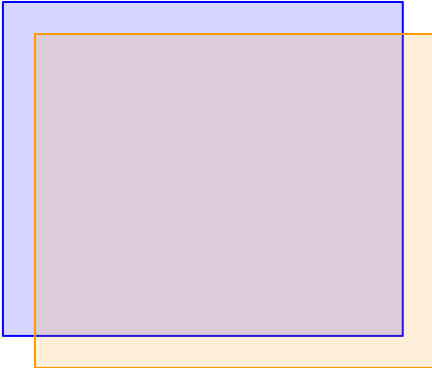
Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(2)} & w_{12}^{(2)} \\ w_{21}^{(2)} & w_{22}^{(2)} \end{bmatrix} =$$

More Parameters


“2x2 Filter”

“2 Channels”





Layer 3



Layer 2



Layer 1

*Convolutional Deep Belief Networks
for Scalable Unsupervised Learning
of Hierarchical Representations, Lee
H., Grosse R., Ranganath R., Ng A.*

The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”

RGB Image

$$\begin{bmatrix} \begin{bmatrix} 1 \\ 0 \\ 1 \\ .5 \end{bmatrix} \begin{bmatrix} .2 \\ .2 \\ 1 \\ 0 \end{bmatrix} \begin{bmatrix} .5 \\ .25 \\ .25 \\ 10 \end{bmatrix} \begin{bmatrix} .15 \\ .25 \\ 1 \\ 0 \end{bmatrix} \begin{bmatrix} 1 \\ .5 \\ 10 \\ .75 \end{bmatrix} \begin{bmatrix} 0 \\ 1 \\ .5 \\ 0 \end{bmatrix} \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix} \begin{bmatrix} 0 \\ 1 \\ 1 \\ 0 \end{bmatrix} \end{bmatrix}$$

The Convolution


- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”

RGB Image 4x4x3

$$\begin{bmatrix} \begin{bmatrix} 1 \\ 0 \\ 1 \\ .5 \end{bmatrix} \begin{bmatrix} .2 \\ .2 \\ 1 \\ 0 \end{bmatrix} \begin{bmatrix} .5 \\ .25 \\ .25 \\ 1 \end{bmatrix} \begin{bmatrix} .15 \\ .25 \\ 1 \\ 0 \end{bmatrix} \begin{bmatrix} 1 \\ .5 \\ 1 \\ .25 \end{bmatrix} \begin{bmatrix} 0 \\ 1 \\ .5 \\ 0 \end{bmatrix} \begin{bmatrix} 0 \\ 0 \\ 1 \\ .2 \end{bmatrix} \end{bmatrix}$$


The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”




The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”




The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”
- Filter channels must match input channels!!!




The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”
- Filter channels must match input channels!!!




The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”
- Filter channels must match input channels!!!




The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”
- Filter channels must match input channels!!!




The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”
- Filter channels must match input channels!!!



The Convolution



Convolution Hyperparameters

- Number of Filters

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”

Stride 1

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”


Stride 1

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”

Stride 1



Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”

Stride 1

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”


Stride 2

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”

Stride 2



Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”

Stride 2

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - What is the dimension of the output for Stride 1 vs. Stride 2?

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
- Size of filter

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 5 & 1 \\ 0 & 1 & 2 \\ 1 & 1 & 0 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
- Size of filter
 - What is output dimension here if stride = 1?

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 5 & 1 \\ 0 & 1 & 2 \\ 1 & 1 & 0 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
- Size of filter
- Problem: size of output keep shrinking!
 - Only a few convolutional layers before the resulting 2D dimensions are very small

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
- Size of filter
- Problem: size of output keep shrinking!
 - Only a few convolutional layers before the resulting 2D dimensions are very small
- Solution: Zero padding

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & .5 & 1 & 0 & 0 \\ 0 & 0 & .25 & .5 & 1 & 0 \\ 0 & 1 & .25 & 0 & 1 & 0 \\ 0 & .5 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Convolution Hyperparameters

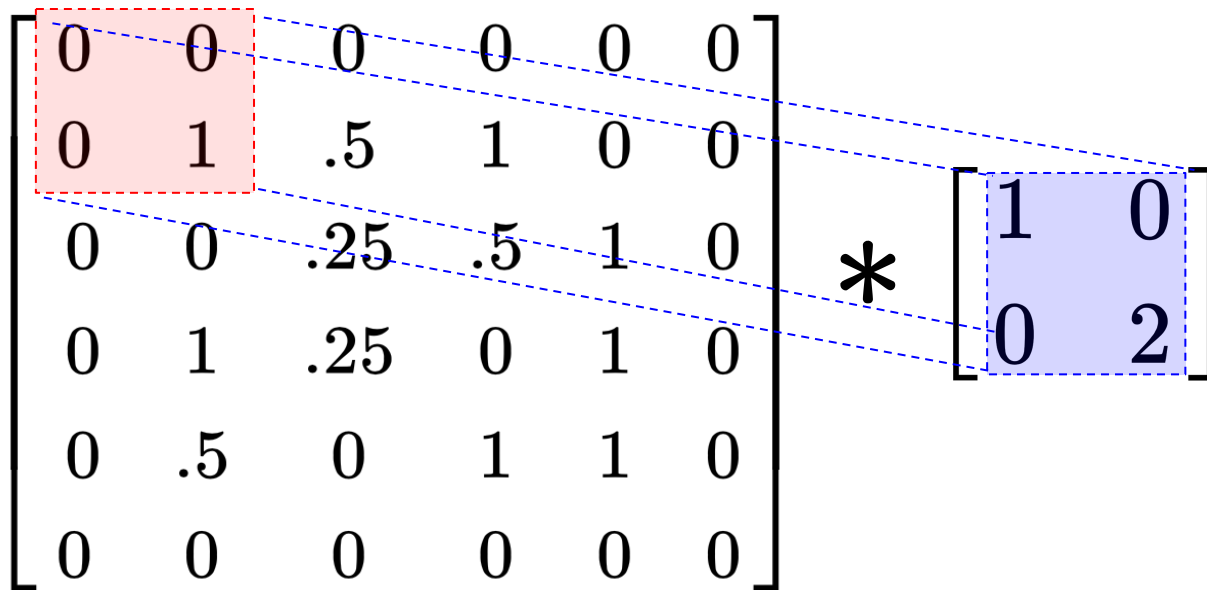
- Number of Filters
- Stride of the filter
- Size of filter
- Padding

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & .5 & 1 & 0 & 0 \\ 0 & 0 & .25 & .5 & 1 & 0 \\ 0 & 1 & .25 & 0 & 1 & 0 \\ 0 & .5 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
- Size of filter
- Padding

Padding
by one



The diagram illustrates a 1D convolution operation. On the left, a 6x6 input matrix is shown with a red dashed box highlighting the top-left 2x2 region (values 0, 0, 0, 1). This region is connected by blue dashed lines to a 2x2 kernel matrix on the right (values 1, 0, 0, 2), which is highlighted with a blue dashed box. A black asterisk (*) is placed between the two matrices, indicating multiplication. An arrow points from the text 'Padding by one' to the top-right corner of the input matrix, where the padding is applied. The input matrix is:

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & .5 & 1 & 0 & 0 \\ 0 & 0 & .25 & .5 & 1 & 0 \\ 0 & 1 & .25 & 0 & 1 & 0 \\ 0 & .5 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

The kernel matrix is:

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Common choices for a Conv-Layer:
 - Stride = 1
 - Odd Filter Size (3x3, 5x5, etc.)
 - “Same” padding

The diagram illustrates a 1D convolution operation. The input is a 6x6 matrix (with padding) and the output is a 4x4 matrix. A 3x3 region of the input is highlighted in red, and a 3x3 region of the output is highlighted in blue. Dashed blue lines show the mapping from the input region to the output region.

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & .5 & 1 & 0 & 0 \\ 0 & 0 & .25 & .5 & 1 & 0 \\ 0 & 1 & .25 & 0 & 1 & 0 \\ 0 & .5 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} * \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & .5 & 2 \end{bmatrix}$$

Convolution Hyperparameters


- Common choices for a Conv-Layer:
 - Stride = 1
 - Odd Filter Size (3x3, 5x5, etc.)
 - “Same” padding

The diagram illustrates a 1D convolution operation with 'Same' padding. The input is a 6x6 matrix, and the kernel is a 3x3 matrix. The output is a 6x6 matrix. The top-left 3x3 region of the input is highlighted in red, and the kernel is highlighted in blue. Blue dashed lines show the mapping from the red region to the kernel and from the kernel to the top-left 3x3 region of the output.

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & .5 & 1 & 0 & 0 \\ 0 & 0 & .25 & .5 & 1 & 0 \\ 0 & 1 & .25 & 0 & 1 & 0 \\ 0 & .5 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} * \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & .5 & 2 \end{bmatrix}$$


Convolutional Neural Network

- LeNet-5 by Yann LeCun




Convolutional Neural Network

- LeNet-5 by Yann LeCun




Convolutional Neural Network

- LeNet-5 by Yann LeCun




Convolutional Neural Network

- LeNet-5 by Yann LeCun




Convolutional Neural Network

- LeNet-5 by Yann LeCun




Convolutional Neural Network

- LeNet-5 by Yann LeCun



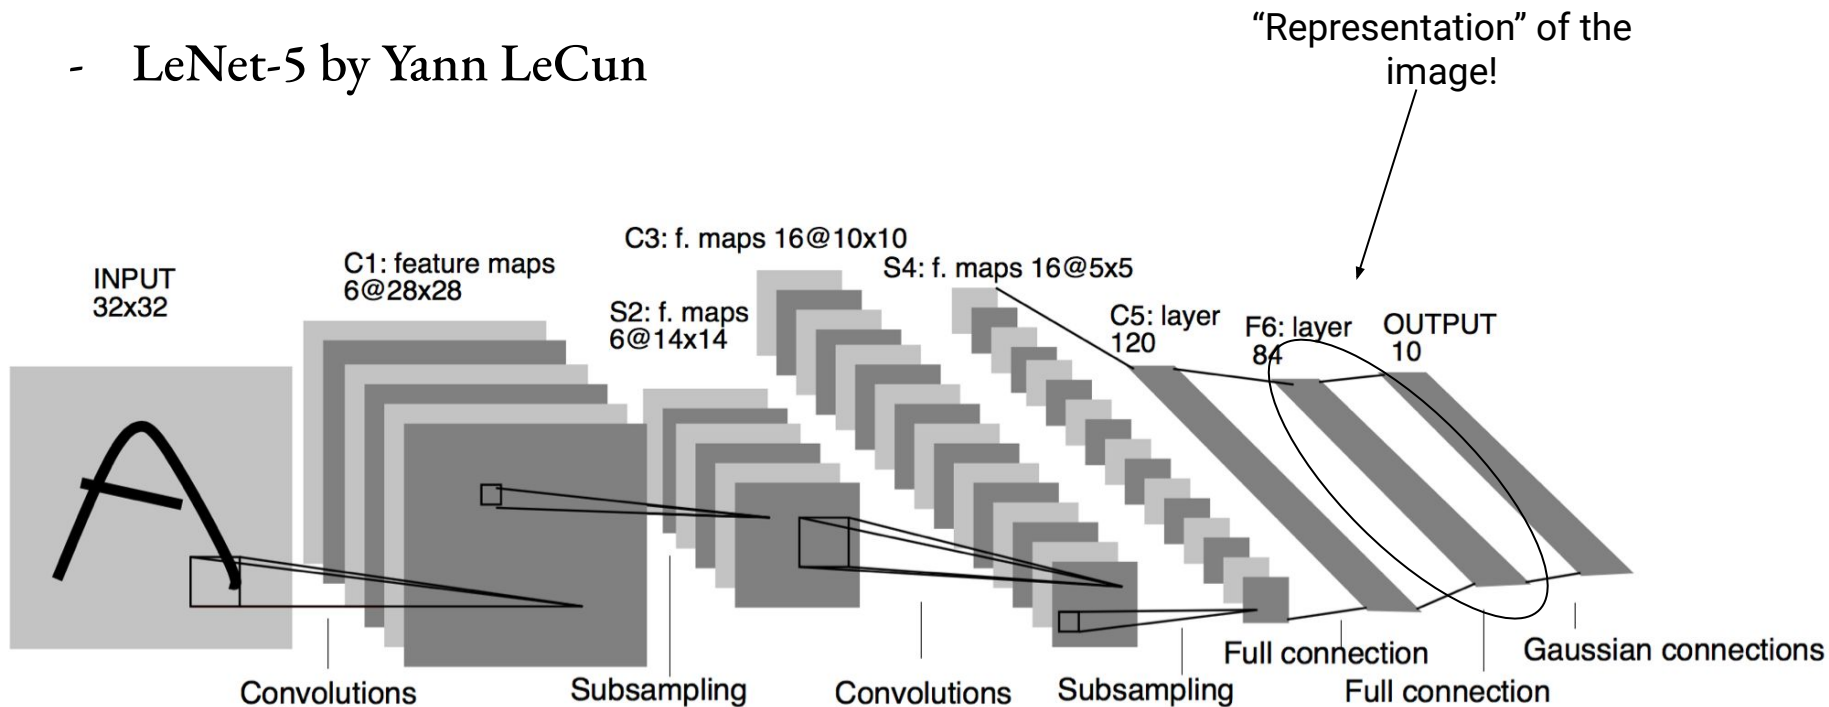
Convolutional Neural Network

- LeNet-5 by Yann LeCun



Convolutional Neural Network

- LeNet-5 by Yann LeCun



Convolutional Neural Network

- AlexNet wins ImageNet Competition in 2012




Figure 2: An illustration of the architecture of our CNN, explicitly showing the delineation of responsibilities between the two GPUs. One GPU runs the layer-parts at the top of the figure while the other runs the layer-parts at the bottom. The GPUs communicate only at certain layers. The network's input is 150,528-dimensional, and the number of neurons in the network's remaining layers is given by 253,440–186,624–64,896–64,896–43,264–4096–4096–1000.

Convolutional Neural Network

- AlexNet wins ImageNet Competition in 2012
- By 2015 we have CNNs with >100 layers, better than human-level performance





Figure 2: An illustration of the architecture of our CNN, explicitly showing the delineation of responsibilities between the two GPUs. One GPU runs the layer-parts at the top of the figure while the other runs the layer-parts at the bottom. The GPUs communicate only at certain layers. The network's input is 150,528-dimensional, and the number of neurons in the network's remaining layers is given by 253,440–186,624–64,896–64,896–43,264–4096–4096–1000.

Downsampling

- Reduce size of output
- Minimal information loss in practice
- Intuition: reduce resolution of the image


Downsampling

- Reduce size of output
- Minimal information loss in practice
- Intuition: reduce resolution of the image
- Max Pooling




Downsampling

- Reduce size of output
 - Minimal information loss in practice
 - Intuition: reduce resolution of the image
 - Max Pooling
- 2x2 filter size
 - Stride 2




Downsampling

- Reduce size of output
 - Minimal information loss in practice
 - Intuition: reduce resolution of the image
 - Max Pooling
- 2x2 filter size
 - Stride 2




Downsampling

- Reduce size of output
 - Minimal information loss in practice
 - Intuition: reduce resolution of the image
 - Max Pooling
 - Average Pooling
- 2x2 filter size
 - Stride 2




Downsampling

- Done along spatial dimension, preserves channels



Convolutional Neural Network

- LeNet-5 by Yann LeCun




Summary

- Convolution Layers
 - Suited for Spatial Data
 - Less Parameters than FC Layers, Weight sharing
- Common Hyperparameters
 - Number of Filters, Filter Size, Stride, Padding
- Common Sequence
 - Conv -> Activation -> Conv -> Activation -> Downsampling
 - Repeat until unrolled into final FC layers


Deeper NNs

- After the success of AlexNet, CNNs got deeper
- Why not just start with as many layers as possible?




Deeper NNs

- After the success of AlexNet, CNNs got deeper
- Why not just start with as many layers as possible?
 - Computer power, data




Deeper NNs

- After the success of AlexNet, CNNs got deeper
- Why not just start with as many layers as possible?
 - Computer power, data
 - Problems with training (vanishing/exploding gradients)



Vanishing/Exploding Gradients




$$a^{[1]} = h(W^{[1]}x + b^{[1]})$$

$$a^{[2]} = h(W^{[2]}a^{[1]} + b^{[2]})$$

$$F(x; \theta) = h(W^{[3]}a^{[2]} + b^{[3]})$$

Vanishing/Exploding Gradients



$$a^{[1]} = h(W^{[1]}x + b^{[1]})$$

$$a^{[2]} = h(W^{[2]}a^{[1]} + b^{[2]})$$

$$F(x; \theta) = h(W^{[3]}a^{[2]} + b^{[3]})$$

$$F = f_1(w_1, f_2(w_2, f_3(w_3)))$$

Vanishing/Exploding Gradients

$$F = f_1(w_1, f_2(w_2, f_3(w_3)))$$

Vanishing/Exploding Gradients

$$F = f_1(w_1, f_2(w_2, f_3(w_3)))$$

$$\frac{\partial F}{\partial w_1} = \frac{\partial f_1}{\partial w_1}$$

Vanishing/Exploding Gradients

$$F = f_1(w_1, f_2(w_2, f_3(w_3)))$$

$$\frac{\partial F}{\partial w_1} = \frac{\partial f_1}{\partial w_1}$$
$$\frac{\partial F}{\partial w_2} = \frac{\partial f_1}{\partial f_2} \frac{\partial f_2}{\partial w_2}$$

Vanishing/Exploding Gradients

$$F = f_1(w_1, f_2(w_2, f_3(w_3)))$$

$$\frac{\partial F}{\partial w_1} = \frac{\partial f_1}{\partial w_1}$$

$$\frac{\partial F}{\partial w_2} = \frac{\partial f_1}{\partial f_2} \frac{\partial f_2}{\partial w_2}$$

$$\frac{\partial F}{\partial w_3} = \frac{\partial f_1}{\partial f_2} \frac{\partial f_2}{\partial f_3} \frac{\partial f_3}{\partial w_3}$$

Vanishing/Exploding Gradients

$$F = f_1(w_1, f_2(w_2, f_3(w_3)))$$

$$(.1)^3 = .001 \quad \frac{\partial F}{\partial w_1} = \frac{\partial f_1}{\partial w_1}$$
$$\frac{\partial F}{\partial w_2} = \frac{\partial f_1}{\partial f_2} \frac{\partial f_2}{\partial w_2}$$
$$\frac{\partial F}{\partial w_3} = \frac{\partial f_1}{\partial f_2} \frac{\partial f_2}{\partial f_3} \frac{\partial f_3}{\partial w_3}$$

Vanishing/Exploding Gradients

$$F = f_1(w_1, f_2(w_2, f_3(w_3)))$$

$$(2)^3 = 8 \qquad \frac{\partial F}{\partial w_1} = \frac{\partial f_1}{\partial w_1}$$

$$\frac{\partial F}{\partial w_2} = \frac{\partial f_1}{\partial f_2} \frac{\partial f_2}{\partial w_2}$$


$$\frac{\partial F}{\partial w_3} = \frac{\partial f_1}{\partial f_2} \frac{\partial f_2}{\partial f_3} \frac{\partial f_3}{\partial w_3}$$

Deeper NNs

- Early parameters can either get stuck, or become unstable during training


Deeper NNs

- Early parameters can either get stuck, or become unstable during training
- Skip Connection




Deeper NNs

- Early parameters can either get stuck, or become unstable during training
- Skip Connection
 - Gradient of earlier parameters depends more directly on output




Deeper NNs

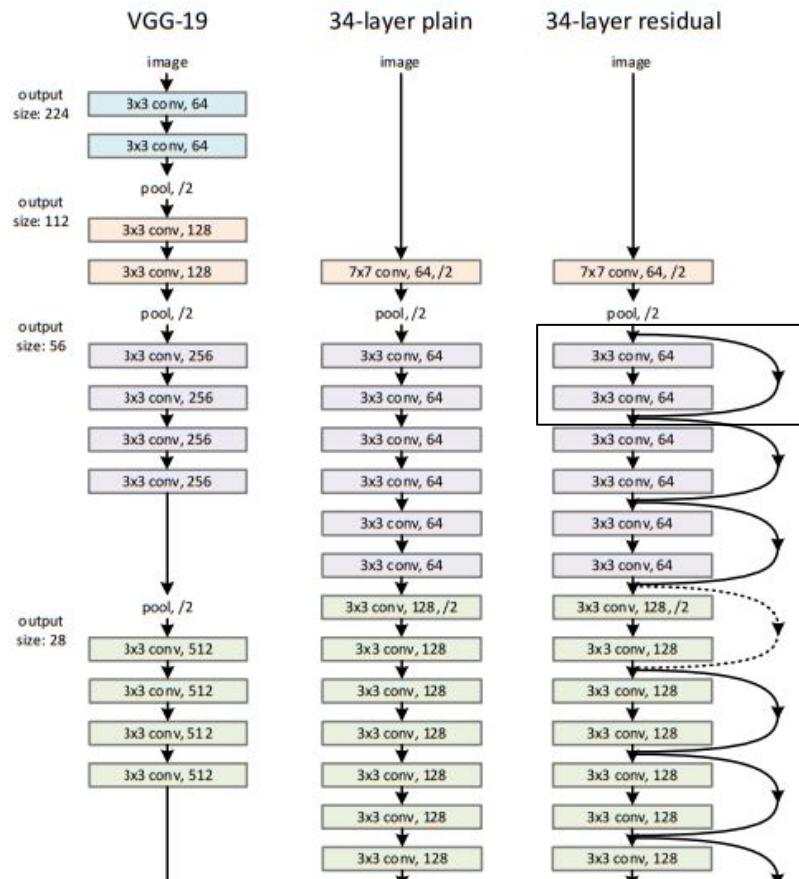
- Early parameters can either get stuck, or become unstable during training
- Skip Connection
 - Gradient of earlier parameters depends more directly on output
 - Identity function easier to learn




VGG vs. ResNet



VGG vs. ResNet



“Residual Block”



Other Techniques

- 1x1 Convolutions
 - With a 1x1 filter size you can condense the channel dimension
- Up-convolution
 - “Up-sample” to increase resolution using parameters
 - UNet
- Adaptive Pooling for Fully Convolutional Networks (FCNs)
 - Pool different shaped images to get same size output
- Normalization
 - Batch Normalization, Layer Normalization, Group Normalization
- 1D/3D Convolutions
 - For 3D: filter size maybe 3x3x3, input is of size (C,H,W,L)

UNet

