

More Imaging Tasks

- Classification
- Objects
 - Object Localization
 - Object Detection
- Keypoint detection
 - Human pose detection
 - Facial recognition
- Segmentation
 - Instance/Scene segmentation
 - Medical imaging

Object Localization

- Where is an object in the image?



Object Localization

- Where is an object in the image?
 - Input: image of a cat
 - Label: location of cat

bounding box



Object Localization

- Where is an object in the image?
 - Input: image of a cat
 - Label: bounding box coordinates

bounding box




- Center/Dimensions: (x, y, w, h)

Object Localization

- Where is an object in the image?
 - Input: image of a cat
 - Label: bounding box coordinates

bounding box

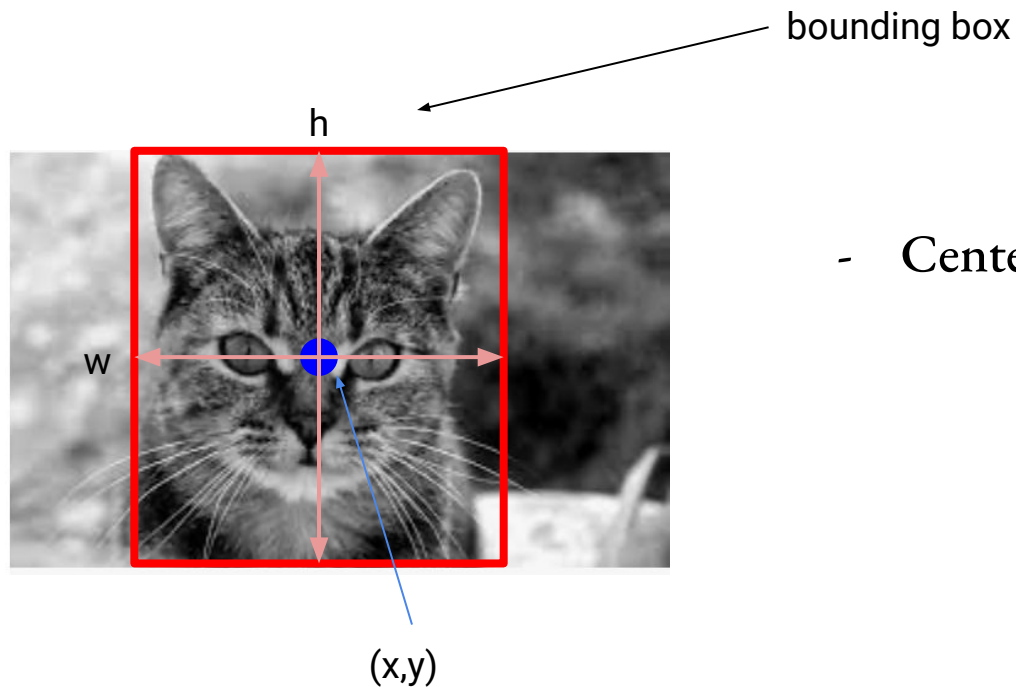


(x,y)

- Center/Dimensions: (x, y, w, h)

Object Localization

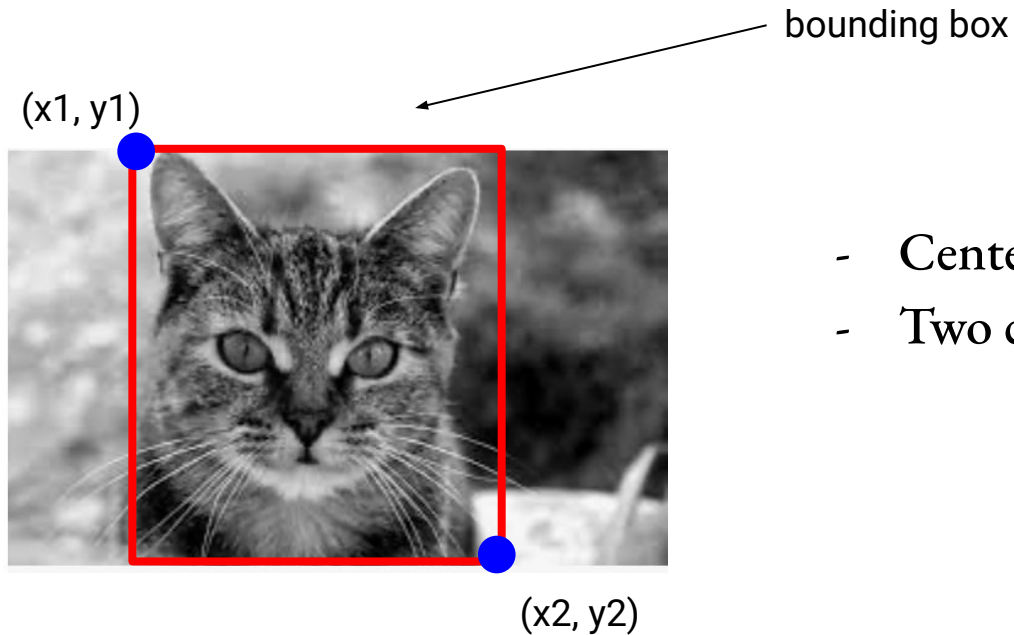
- Where is an object in the image?
 - Input: image of a cat
 - Label: bounding box coordinates



- Center/Dimensions: (x, y, w, h)

Object Localization

- Where is an object in the image?
 - Input: image of a cat
 - Label: bounding box coordinates



- Center/Dimensions: (x, y, w, h)
- Two corners: $(x1, y1, x2, y2)$

Object Localization

- Where is an object in the image?
 - Input: image of a cat
 - Label: bounding box coordinates

bounding box



- Center/Dimensions: (x, y, w, h)
- Two corners: (x_1, y_1, x_2, y_2)

Object Localization

- Where is an object in the image?
 - Input: image of a cat
 - Label: bounding box coordinates



bounding box

regression

- Center/Dimensions: (x, y, w, h)
- Two corners: (x_1, y_1, x_2, y_2)

Classification + Localization

- Simultaneously answer:
 - Does this photo contain a cat?
 - If so where is the cat?



cat

Classification + Localization

- Simultaneously answer:
 - Does this photo contain a cat?
 - If so where is the cat?



cat

(1, 150, 175, 50, 70) \leftrightarrow (class, x, y, w, h)

Classification + Localization

- Simultaneously answer:
 - Does this photo contain a cat?
 - If so where is the cat?



cat

Classification: BCELoss(predicted class, true class)

Localization: MSELoss(predicted coords, true coords)

(1, 150, 175, 50, 70) \leftrightarrow (class, x, y, w, h)

Classification + Localization


- Simultaneously answer:
 - Does this photo contain a cat?
 - If so where is the cat?



cat

(1, 150, 175, 50, 70) \leftrightarrow (class, x, y, w, h)

Class label



$$Loss = BCELoss + Y * MSELoss$$


Classification + Localization

- Simultaneously answer:
 - Does this photo contain a cat?
 - If so where is the cat?



cat

Class label




$$Loss = BCELoss + \gamma(Y * MSELoss)$$

(1, 150, 175, 50, 70) \leftrightarrow (class, x, y, w, h)

Intersection over Union for Localization

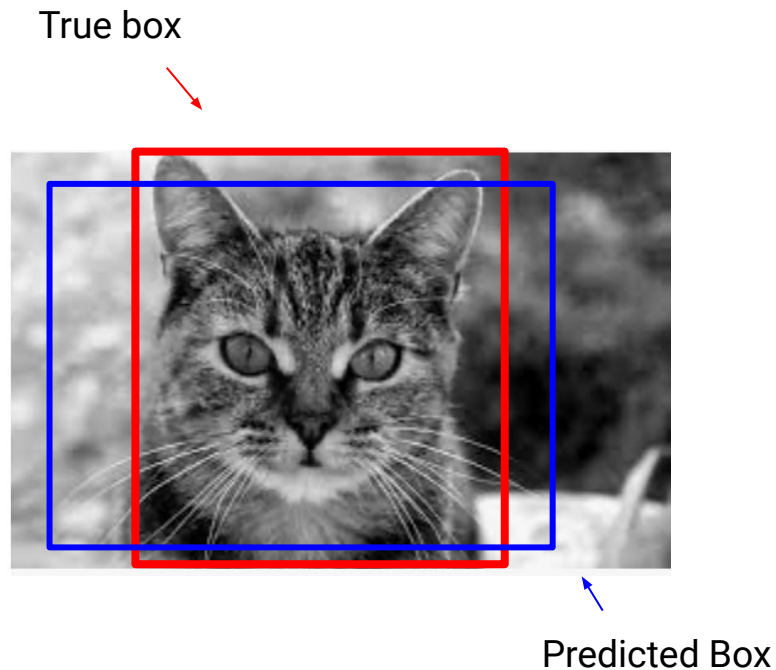
True box



Predicted Box




Intersection over Union for Localization



$$\text{IoU: } \frac{\textit{Intersection Area}}{\textit{Union Area}}$$

- Score from 0 to 1
- Higher is better


Intersection over Union for Localization



$$\text{IoU: } \frac{\textit{Intersection Area}}{\textit{Union Area}}$$

- Score from 0 to 1
- Higher is better


Intersection over Union for Localization



$$\text{IoU: } \frac{\textit{Intersection Area}}{\textit{Union Area}}$$

- Score from 0 to 1
- Higher is better

Intersection over Union for Localization



$$\text{IoU: } \frac{\textit{Intersection Area}}{\textit{Union Area}}$$

- Score from 0 to 1
- Higher is better
- Areas can be computed from box coordinates
- Union given by

$$\text{Box1 Area} + \text{Box2 Area} - \text{Intersection Area}$$

Multi-Label Classification + Localization

- What about multiple classes in the same photo?

Multi-Label Classification + Localization

- What about multiple classes in the same photo?
- k classes \rightarrow k binary classification tasks + localization

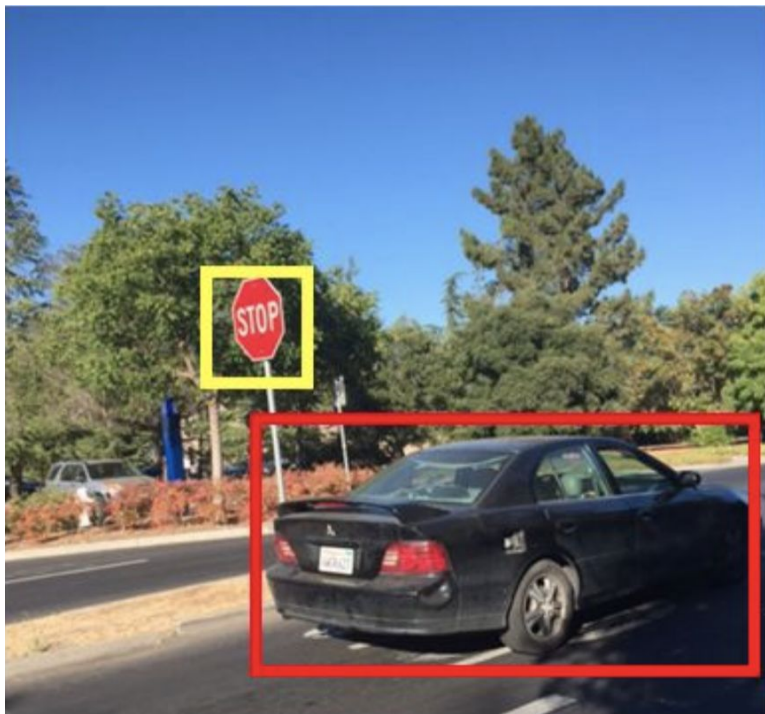


Classes

- Pedestrian
- Stop sign
- Car
- Traffic Light

Multi-Label Classification + Localization

- What about multiple classes in the same photo?
- k classes \rightarrow k binary classification tasks + localization



Classes

- Pedestrian
- Stop sign
- Car
- Traffic Light

Label



(0, 0, 0, 0, 0)

(1, 200, 300, 10, 10)

(1, 400, 100, 300, 100)

(0, 0, 0, 0, 0)

Multi-Label Classification + Localization

- What about multiple classes in the same photo?
- k classes \rightarrow k binary classification tasks + localization



Classes

- Pedestrian
- Stop sign
- Car
- Traffic Light

Label



classes

(0, 0, 0, 0, 0)
(1, 200, 300, 10, 10)
(1, 400, 100, 300, 100)
(0, 0, 0, 0, 0)

Multi-Label Classification + Localization

- What about multiple classes in the same photo?
- k classes -> k binary classification tasks + localization



Classes

- Pedestrian
- Stop sign
- Car
- Traffic Light

Label

Bounding box
coordinates

classes

(0, 0, 0, 0, 0)
(1, 200, 300, 10, 10)
(1, 400, 100, 300, 100)
(0, 0, 0, 0, 0)

Object Detection


- Classify and locate all objects in an image

Object Detection

- Classify and locate all objects in an image
- Region Proposals (RCNNs)
 1. Propose region for object classification + localization
 2. Perform object classification + localization on this region


Object Detection

- Classify and locate all objects in an image
- Region Proposals (RCNNs)
 1. Propose region for object classification + localization
 2. Perform object classification + localization on this region




Object Detection

- Classify and locate all objects in an image
- Region Proposals (RCNNs)
 1. Propose region for object classification + localization
 2. Perform object classification + localization on this region






Object Detection

- Classify and locate all objects in an image
- Region Proposals (RCNNs)
 1. Propose region for object classification + localization
 2. Perform object classification + localization on this region



Object Detection

- Classify and locate all objects in an image
- Region Proposals (RCNNs)
 1. Propose region for object classification + localization
 2. Perform object classification + localization on this region




CNN for Classification + Localization

Object Detection

- Classify and locate all objects in an image
- Region Proposals (RCNNs)
- You Only Look Once (YOLO)
 1. Sort each bounding box label into various “anchor boxes”
 2. Divide Image into a grid
 3. For each grid square predict whether the center of a bounding box appears in the grid square (for each anchor box)



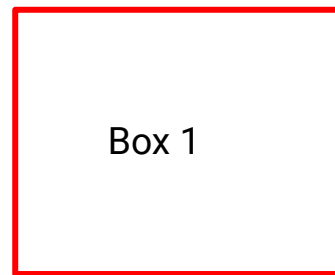

YOLO

1. Sort each bounding box label into various “anchor boxes”
2. Divide Image into a grid
3. For each grid square predict whether the center of a bounding box appears in the grid square (for each anchor box)




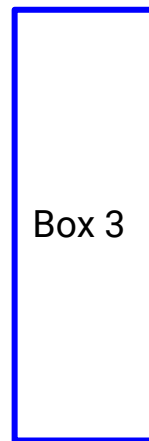
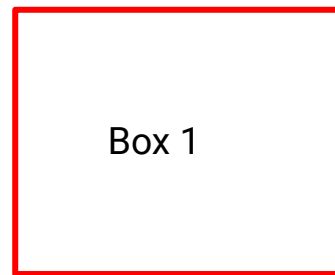
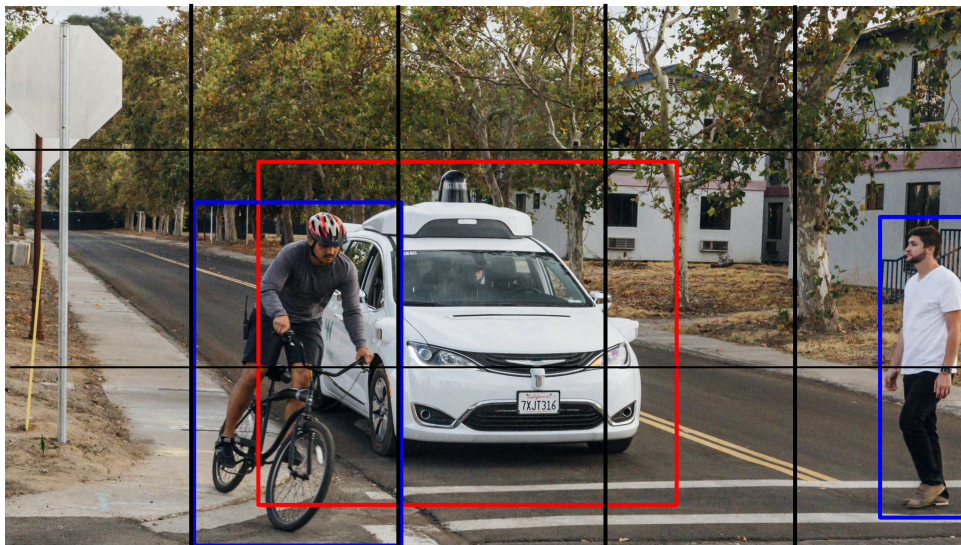
YOLO

1. Sort each bounding box label into various “anchor boxes”
2. Divide Image into a grid
3. For each grid square predict whether the center of a bounding box appears in the grid square (for each anchor box)







YOLO

1. Sort each bounding box label into various “anchor boxes”
2. Divide Image into a grid
3. For each grid square predict whether the center of a bounding box appears in the grid square (for each anchor box)

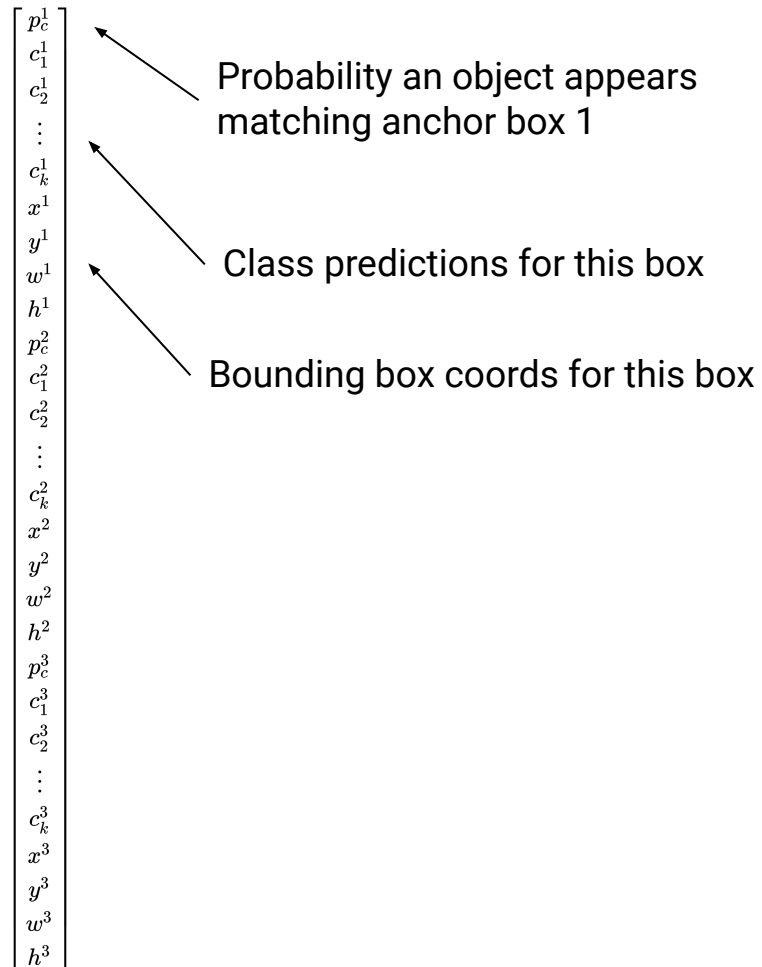



YOLO

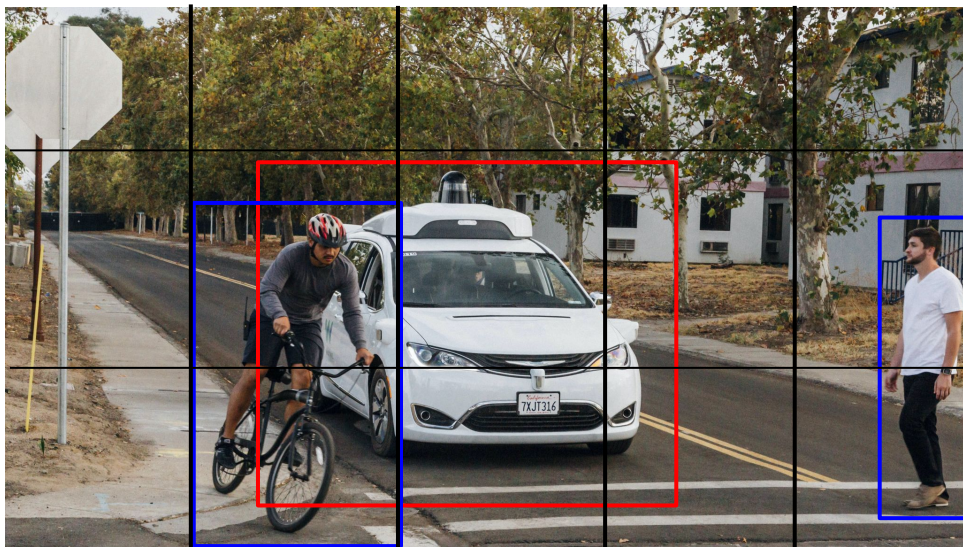
1. Sort each bounding box label into various “anchor boxes”
2. Divide Image into a grid
3. For each grid square predict whether the center of a bounding box appears in the grid square (for each anchor box)



YOLO



YOLO



$$\begin{bmatrix} p_c^1 \\ c_1^1 \\ c_2^1 \\ \vdots \\ c_k^1 \\ x^1 \\ y^1 \\ w^1 \\ h^1 \\ p_c^2 \\ c_1^2 \\ c_2^2 \\ \vdots \\ c_k^2 \\ x^2 \\ y^2 \\ w^2 \\ h^2 \\ p_c^3 \\ c_1^3 \\ c_2^3 \\ \vdots \\ c_k^3 \\ x^3 \\ y^3 \\ w^3 \\ h^3 \end{bmatrix}$$

Probability an object appears
matching anchor box 2

Class predictions for this box

Bounding box coords for this box



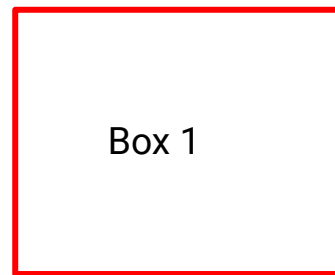
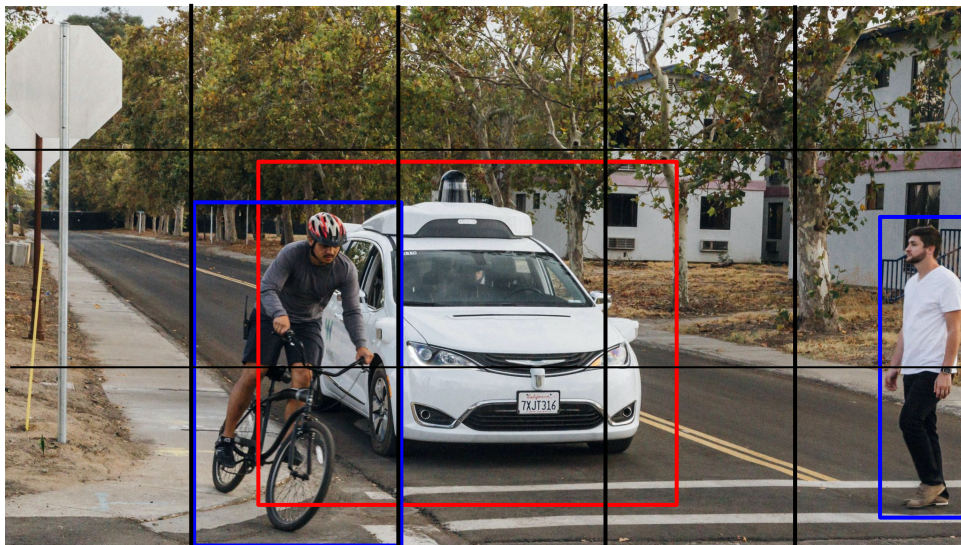
YOLO







$$\begin{bmatrix} (p_c^1)^1 & \dots & (p_c^1)^N \\ (c_1^1)^1 & & (c_1^1)^N \\ (c_2^1)^1 & & (c_2^1)^N \\ \vdots & & \vdots \\ (c_k^1)^1 & & (c_k^1)^N \\ (x^1)^1 & & (x^1)^N \\ (y^1)^1 & & (y^1)^N \\ (w^1)^1 & & (w^1)^N \\ (h^1)^1 & & (h^1)^N \\ (p_c^2)^1 & & (p_c^2)^N \\ (c_1^2)^1 & & (c_1^2)^N \\ (c_2^2)^1 & & (c_2^2)^N \\ \vdots & & \vdots \\ (c_k^2)^1 & & (c_k^2)^N \\ (x^2)^1 & & (x^2)^N \\ (y^2)^1 & & (y^2)^N \\ (w^2)^1 & & (w^2)^N \\ (h^2)^1 & & (h^2)^N \\ (p_c^3)^1 & & (p_c^3)^N \\ (c_1^3)^1 & & (c_1^3)^N \\ (c_2^3)^1 & & (c_2^3)^N \\ \vdots & & \vdots \\ (c_k^3)^1 & & (c_k^3)^N \\ (x^3)^1 & & (x^3)^N \\ (y^3)^1 & & (y^3)^N \\ (w^3)^1 & & (w^3)^N \\ (h^3)^1 & \dots & (h^3)^N \end{bmatrix}$$

YOLO

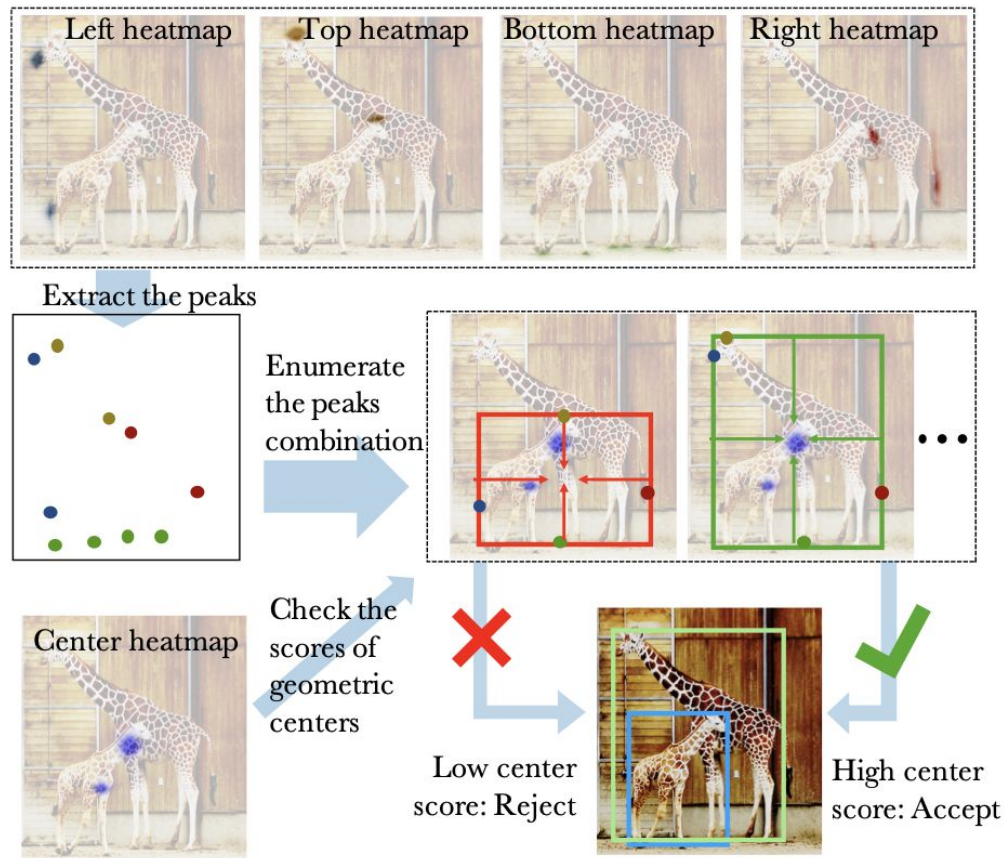
1. Sort each bounding box label into various “anchor boxes”
2. Divide Image into a grid
3. For each grid square predict whether the center of a bounding box appears in the grid square (for each anchor box)





Other Methods

- Some methods output a “heatmap” or series of heatmaps
- Probability each pixel is a keypoint of the bounding box and its class
- Examples
 - CenterNet
 - ExtremeNet



Bottom-up Object Detection by Grouping Extreme and Center Points

Segmentation

- Instance or Scene Segmentation
- More sophisticated labels
- Output/Labels are masks or series of masks (same size as image)




Segmentation

- Medical Imaging




Input





Label

Segmentation

- Medical Imaging




Input



labels


Segmentation

- UNet-like Architecture




Segmentation

- UNet-like Architecture




Segmentation

- UNet-like Architecture



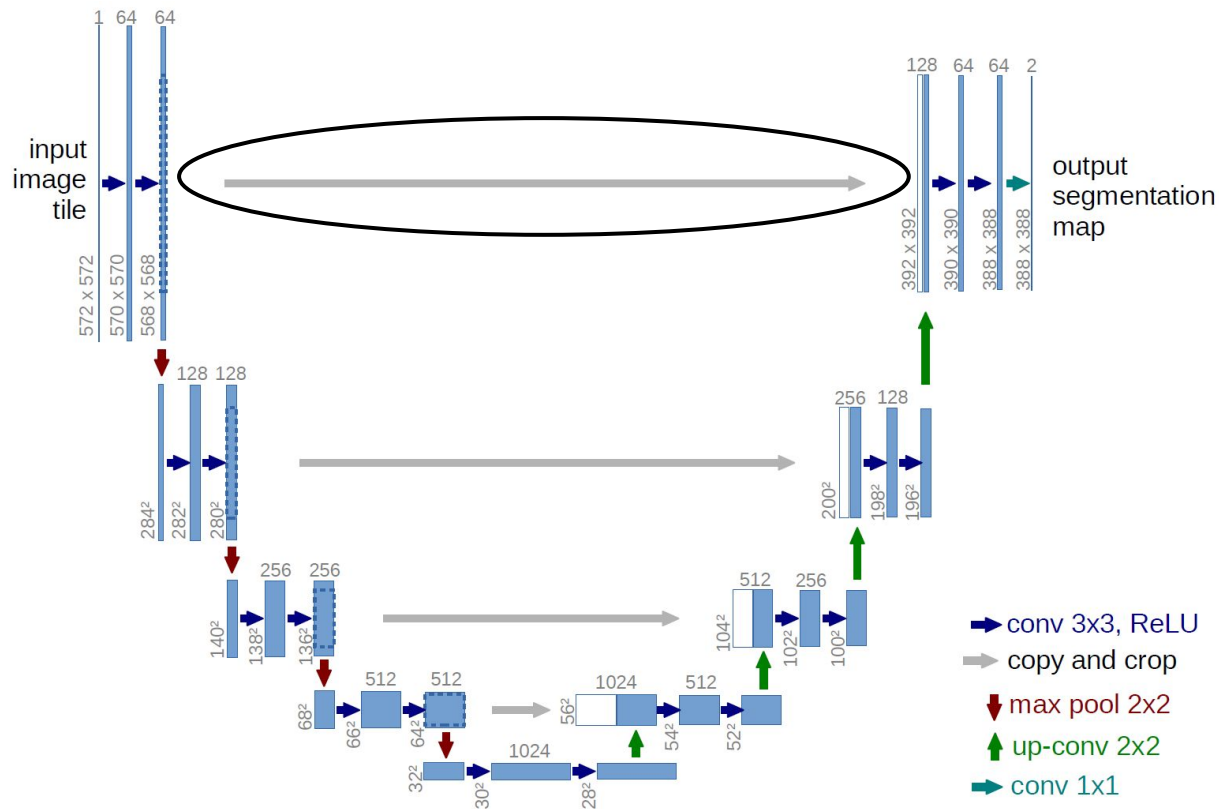
Segmentation

- UNet-like Architecture




Segmentation

- UNet-like Architecture



Up-Sampling

- Increase resolution of the feature map
- Simplest: Un-pooling



Up-Sampling

- Increase resolution of the feature map
- Simplest: Un-pooling
- Up-Convolution


$$\begin{bmatrix} .5 & .25 \\ 0 & 1 \end{bmatrix} * \begin{bmatrix} 0 & -1 \\ 1 & 0 \end{bmatrix}$$

Input

Filter

Up-Sampling

- Increase resolution of the feature map
- Simplest: Un-pooling
- Up-Convolution



Up-Sampling

- Increase resolution of the feature map
- Simplest: Un-pooling
- Up-Convolution

The diagram illustrates a 1D up-convolution operation. It starts with an input vector $\begin{bmatrix} .5 & .25 \\ 0 & 1 \end{bmatrix}$ where the value $.25$ is highlighted in a box. This is multiplied by a filter vector $\begin{bmatrix} 0 & -1 \\ 1 & 0 \end{bmatrix}$. The resulting intermediate vector is $\begin{bmatrix} 0 & -.5 \\ .5 & 0 \end{bmatrix}$. An arrow points from the $.5$ element of this vector to the final output vector $\begin{bmatrix} 0 & -.25 \\ .25 & 0 \end{bmatrix}$, where the $.25$ element is also highlighted in a box. The labels 'Input' and 'Filter' are positioned below their respective vectors.

$$\begin{bmatrix} .5 & .25 \\ 0 & 1 \end{bmatrix} * \begin{bmatrix} 0 & -1 \\ 1 & 0 \end{bmatrix} \rightarrow \begin{bmatrix} 0 & -.5 \\ .5 & 0 \end{bmatrix} \rightarrow \begin{bmatrix} 0 & -.25 \\ .25 & 0 \end{bmatrix}$$

Input Filter

Up-Sampling

- Increase resolution of the feature map
- Simplest: Un-pooling
- Up-Convolution

$$\begin{array}{ccc} \begin{bmatrix} .5 & .25 \\ 0 & 1 \end{bmatrix} & * & \begin{bmatrix} 0 & -1 \\ 1 & 0 \end{bmatrix} \longrightarrow \begin{bmatrix} 0 & -.5 & 0 & -.25 \\ .5 & 0 & .25 & 0 \\ 0 & 0 & 0 & -1 \\ 0 & 0 & 1 & 0 \end{bmatrix} \\ \text{Input} & & \text{Filter} \end{array}$$

Up-Sampling

- Increase resolution of the feature map
- Simplest: Un-pooling
- Up-Convolution

$$\begin{array}{ccc} \begin{bmatrix} .5 & .25 \\ 0 & 1 \end{bmatrix} & * & \begin{bmatrix} w_{11} & w_{12} \\ w_{21} & w_{22} \end{bmatrix} \longrightarrow \begin{bmatrix} 0 & -.5 & 0 & -.25 \\ .5 & 0 & .25 & 0 \\ 0 & 0 & 0 & -1 \\ 0 & 0 & 1 & 0 \end{bmatrix} \\ \text{Input} & & \text{Filter} \end{array}$$