



# Overview

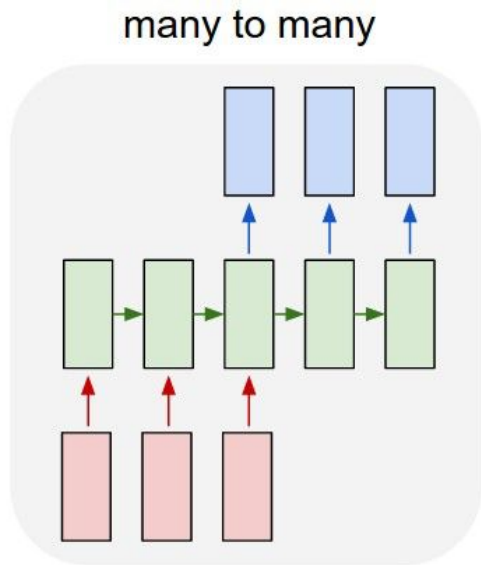

- Attention
- The Transformer
- BERT

# Seq2Seq Models



# Seq2Seq Models


- Information Bottleneck!




# Problems with translation

- Have to encode a full sentence into one vector


Ich muss auf den Markt gehen.



I must go to the Market.



many to many




# Attention


- Use the hidden states generated by the encoder RNN.
- In each decoding step, weight the hidden states (attention mechanism) and use the weighted sum to predict the next element.

# Attention

- Use the hidden states generated by the encoder RNN.
- In each decoding step, weight the hidden states (attention mechanism) and use the weighted sum to predict the next element.



# Attention





# Attention

Context Vector


$$\sum_{i=1}^N \alpha_i^t h_i = c_t$$

$$\sum_{i=1}^N \alpha_i^t = 1$$





# Attention




Context Vector



$c_t$

$$\sum_{i=1}^N \alpha_i^t h_i = c_t$$
$$\sum_{i=1}^N \alpha_i^t = 1$$




# Attention


Context Vector

$$\sum_{i=1}^N \alpha_i^t h_i = c_t$$

$$\sum_{i=1}^N \alpha_i^t = 1$$




# Attention



$$\sum_{i=1}^N \alpha_i^t h_i = c_t$$
$$\sum_{i=1}^N \alpha_i^t = 1$$

# Attention



Context Vector



$$\sum_{i=1}^N \alpha_i^t h_i = c_t$$


$$\sum_{i=1}^N \alpha_i^t = 1$$

Softmax

$$\tilde{\alpha}_i^1 = SCORE(\vec{h}_i, \vec{s}_1)$$

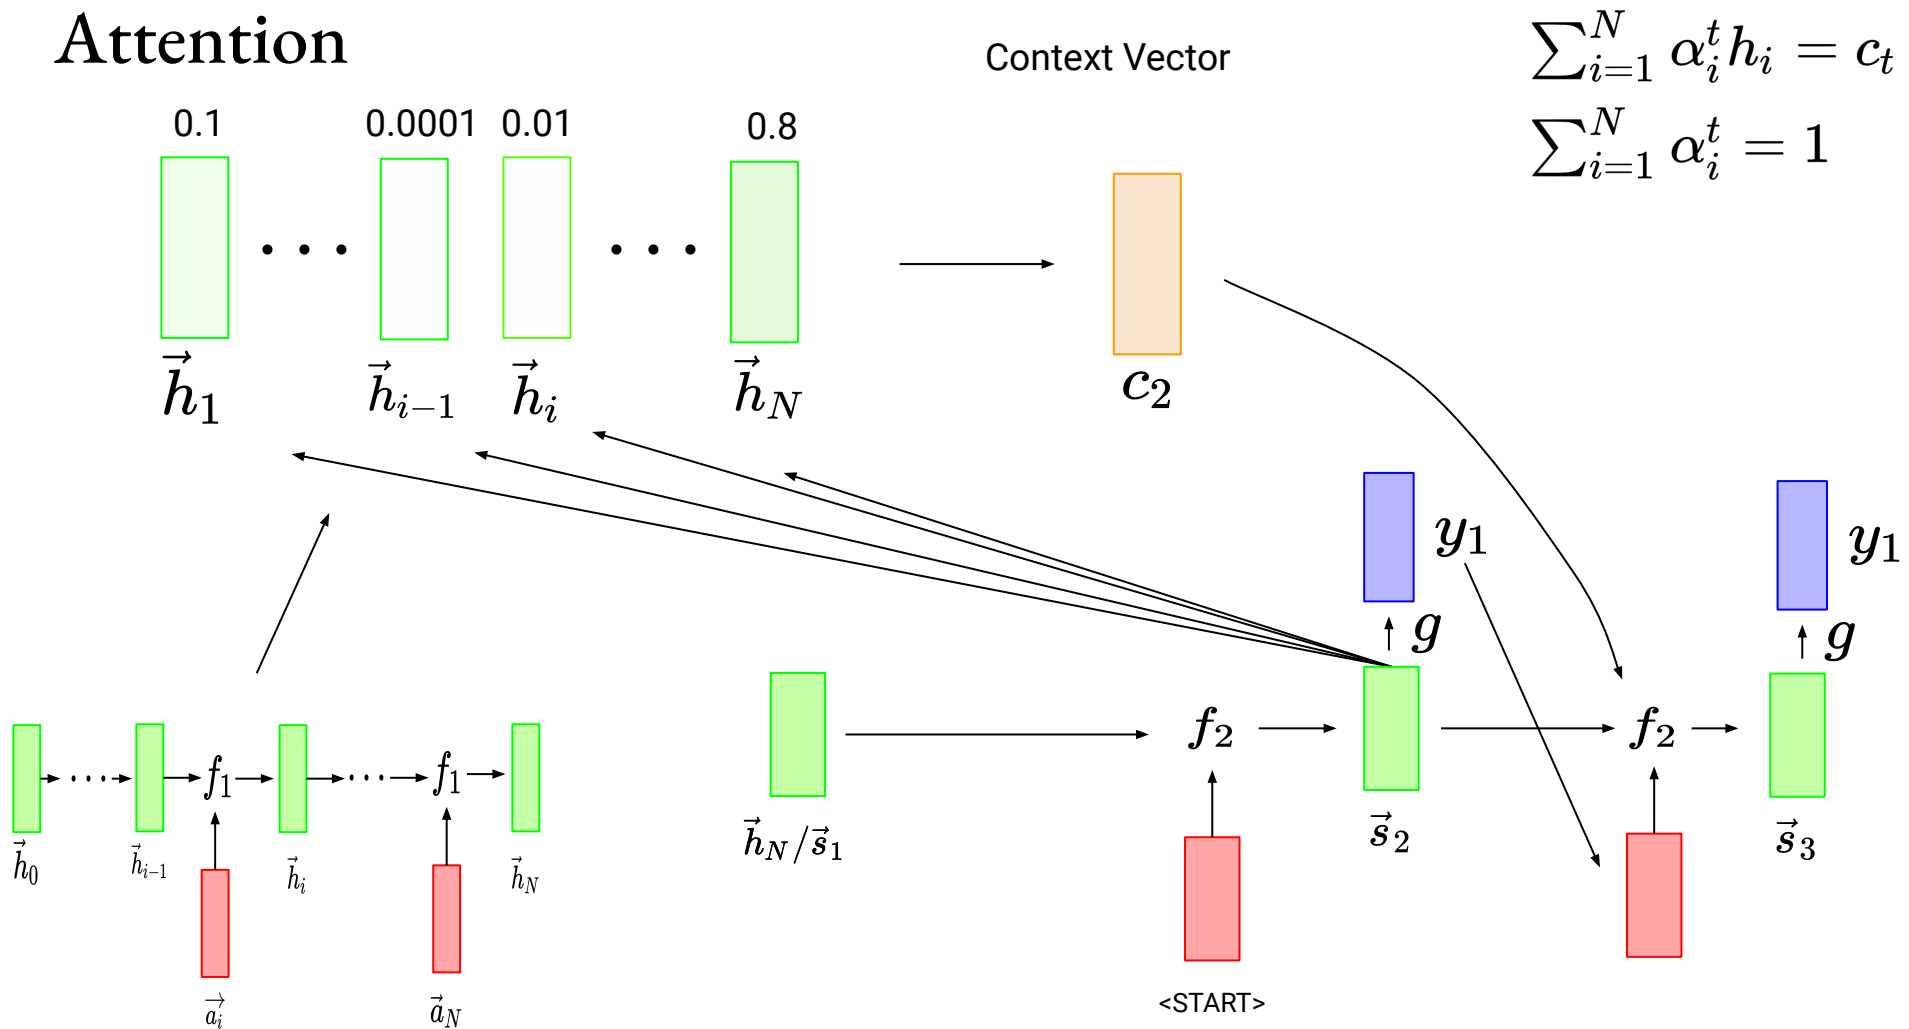
$$\alpha_i^1 = \frac{\exp(\tilde{\alpha}_i^1)}{\sum_{j=1}^N \exp(\tilde{\alpha}_j^1)}$$

# Attention



$$\sum_{i=1}^N \alpha_i^t h_i = c_t$$
$$\sum_{i=1}^N \alpha_i^t = 1$$

# Attention




# Attention

- The Attention Score function can be as simple as:  $\vec{s}_t \cdot \vec{h}_i$
- Or a weight matrix can be introduced:  $\vec{s}_t \cdot W_a \vec{h}_i$ 
  - (More can be found [here](#))

# Attention

- The Attention Score function can be as simple as:  $\vec{s}_t \cdot \vec{h}_i$
- Or a weight matrix can be introduced:  $\vec{s}_t \cdot W_a \vec{h}_i$ 
  - (More can be found [here](#))




Bahdanau et al., 2015

- Attention Matrix
- Introduces some *Explainability*.



# Attention

- The Attention Score function can be as simple as:  $\vec{s}_t \cdot \vec{h}_i$
- Or a weight matrix can be introduced:  $\vec{s}_t \cdot W_a \vec{h}_i$ 
  - (More can be found [here](#))


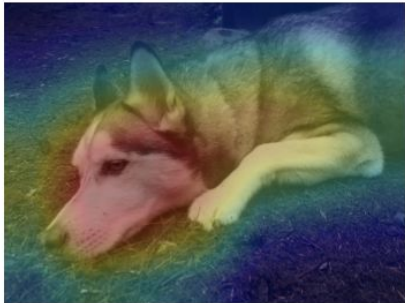



Bahdanau et al., 2015

- Attention Matrix
- Introduces some *Explainability*.



Just because you pay attention to the right thing doesn't mean you are paying attention for the right *reason*.

	Test Image	Evidence for Animal Being a Siberian Husky	Evidence for Animal Being a Transverse Flute
Explanations Using Attention Maps			

*Stop Explaining Black Box Machine Learning Models  
for High Stakes Decisions and Use Interpretable  
Models Instead*  
By Cynthia Rudin 2018  
[link](#)



Just because you pay attention to the right thing doesn't  
mean you are paying attention for the right *reason*.

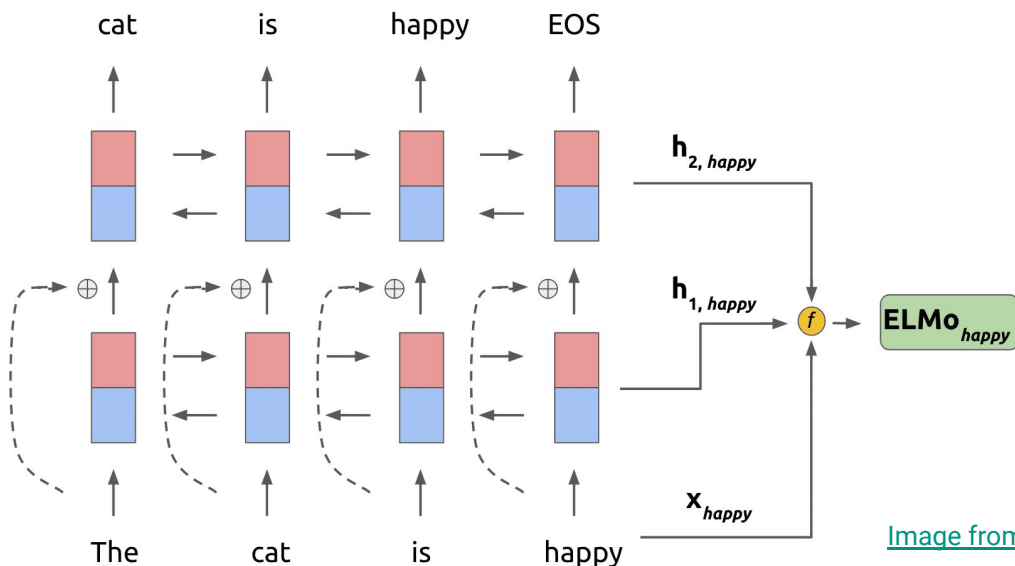
# Contextual Word Embeddings

- What about homonyms?
  - He took a *train*.
  - He likes to *train* in the mornings.

# Contextual Word Embeddings



- ELMo: Embeddings from Language Models
  - Trained on self-supervised task such as predict the next word or Part of Speech tagging
  - Creates contextual word embeddings



Idea: Stack bi-directional LSTMs to produce multiple embeddings of words that use the context.

- (Can get more complicated, but that's the core)

[Image from here](#)


# Self-Attention

- Goal: Only use Attention mechanism to encode a sequence of tokens.
  - Generate contextual embeddings of the tokens without RNNs!

# Self-Attention

$$W_Q, W_K, W_V$$


- Initial Embeddings
- 3 weight matrices: Query, Key Value



# Self-Attention

$$W_Q, W_K, W_V$$


- Each token gets a Query, Key, and Value vector.



# Self-Attention

$$W_Q, W_K, W_V$$

- For each vector we use *its query* to explore *all the keys* which generate weights.
- The resulting embedding is the weighted sum of all the value vectors.






# Self-Attention

$$W_Q, W_K, W_V$$

- For each vector we use *its query* to explore *all the keys* which generate weights.
- The resulting embedding is the weighted sum of all the value vectors.




- We first take the dot product of  $q_1$  with every key vector.
- We get the weights by using a softmax layer.
- Sometimes there is a constant multiplier in there as well (for stability purposes)

# Self-Attention

$$W_Q, W_K, W_V$$

- For each vector we use *its query* to explore *all the keys* which generate weights.
- The resulting embedding is the weighted sum of all the value vectors.




- We first take the dot product of  $q_1$  with every key vector.
- We get the weights by using a softmax layer.
- Sometimes there is a constant multiplier in there as well (for stability purposes)

$$\sum_{i=1}^N \alpha_i^t \vec{v}_i$$

# Self-Attention

$$W_Q, W_K, W_V$$

- For each vector we use *its query* to explore *all the keys* which generate weights.
- The resulting embedding is the weighted sum of all the value vectors.




- We first take the dot product of  $q_1$  with every key vector.
- We get the weights by using a softmax layer.
- Sometimes there is a constant multiplier in there as well (for stability purposes)

$$\sum_{i=1}^N \alpha_i^t \vec{v}_i$$



# Self-Attention

- Contextual Embeddings!

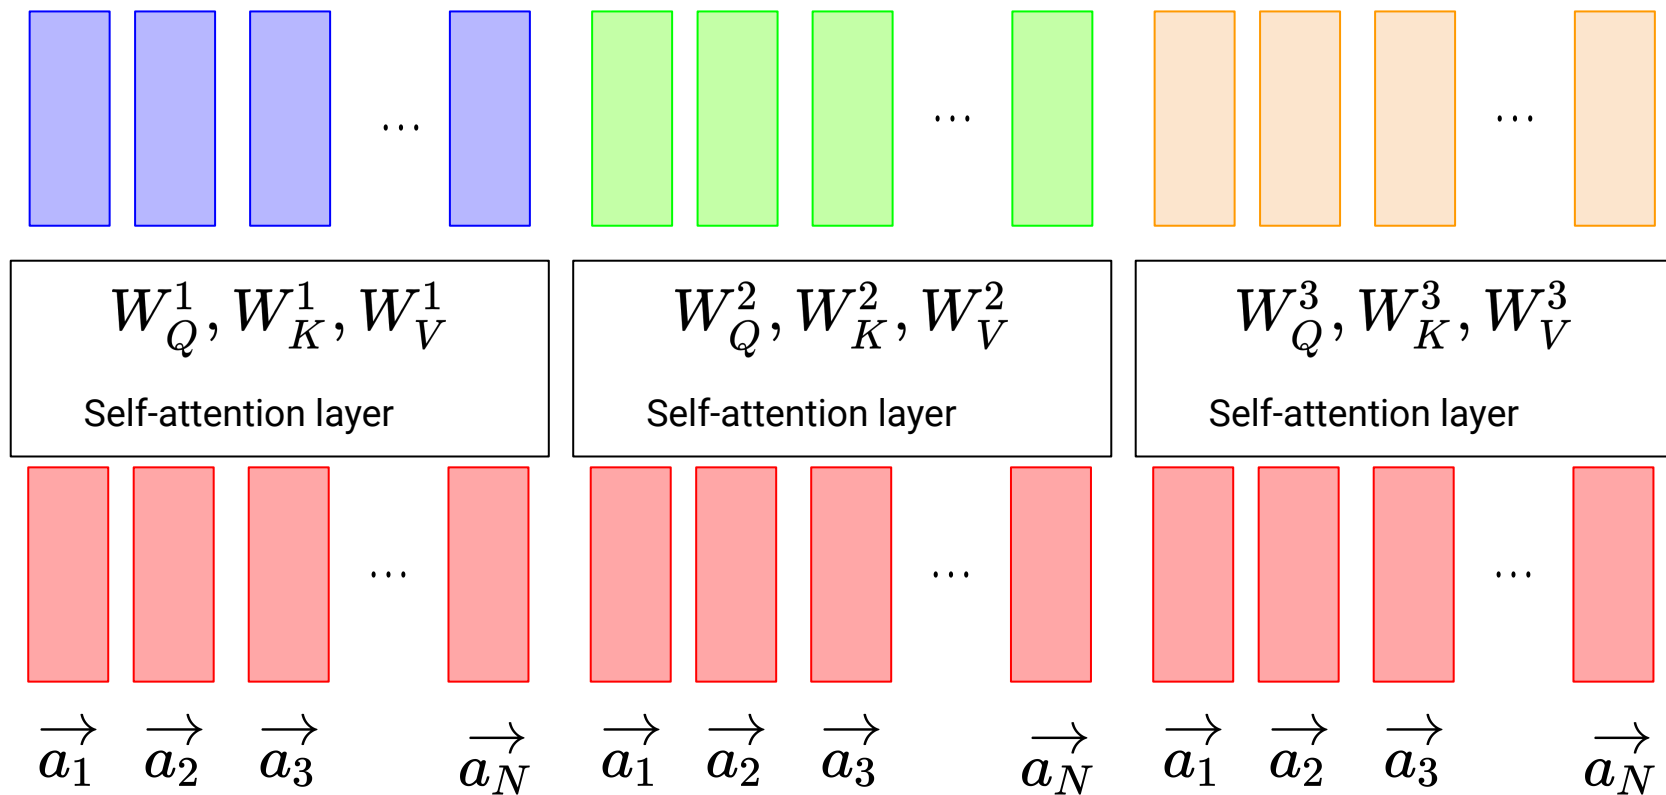


# Self-Attention

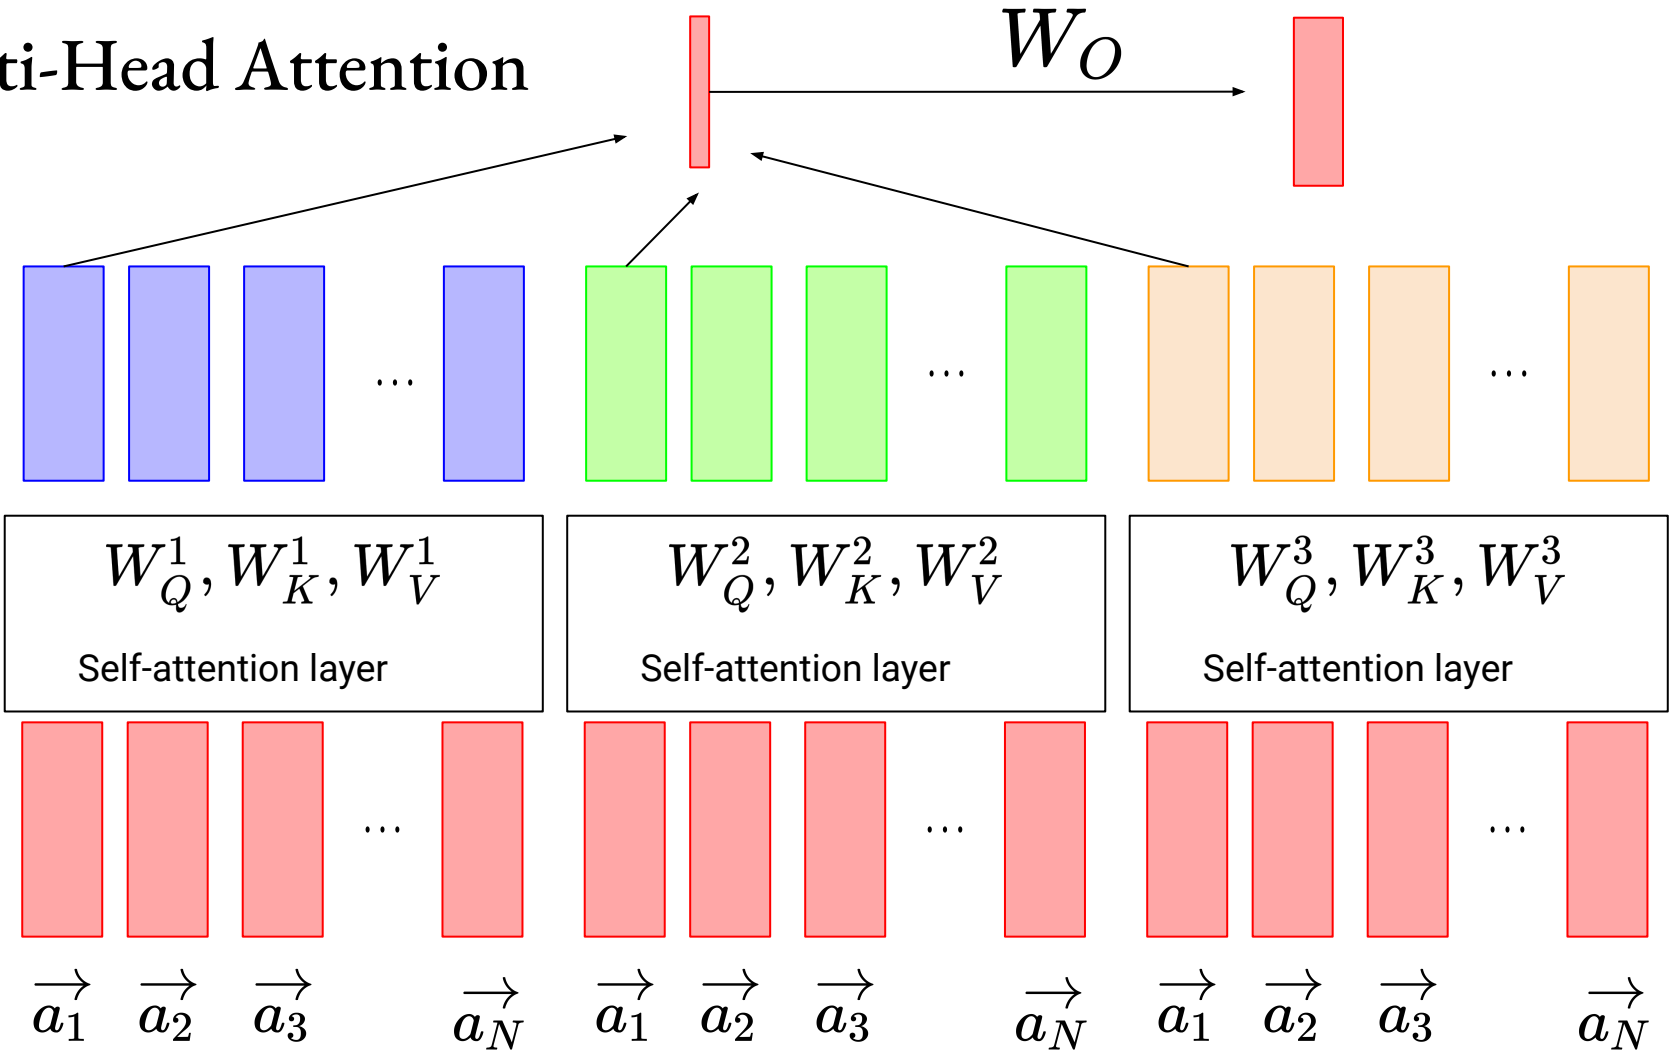
- Sequence Agnostic! Often Positional Embeddings are added here as well.



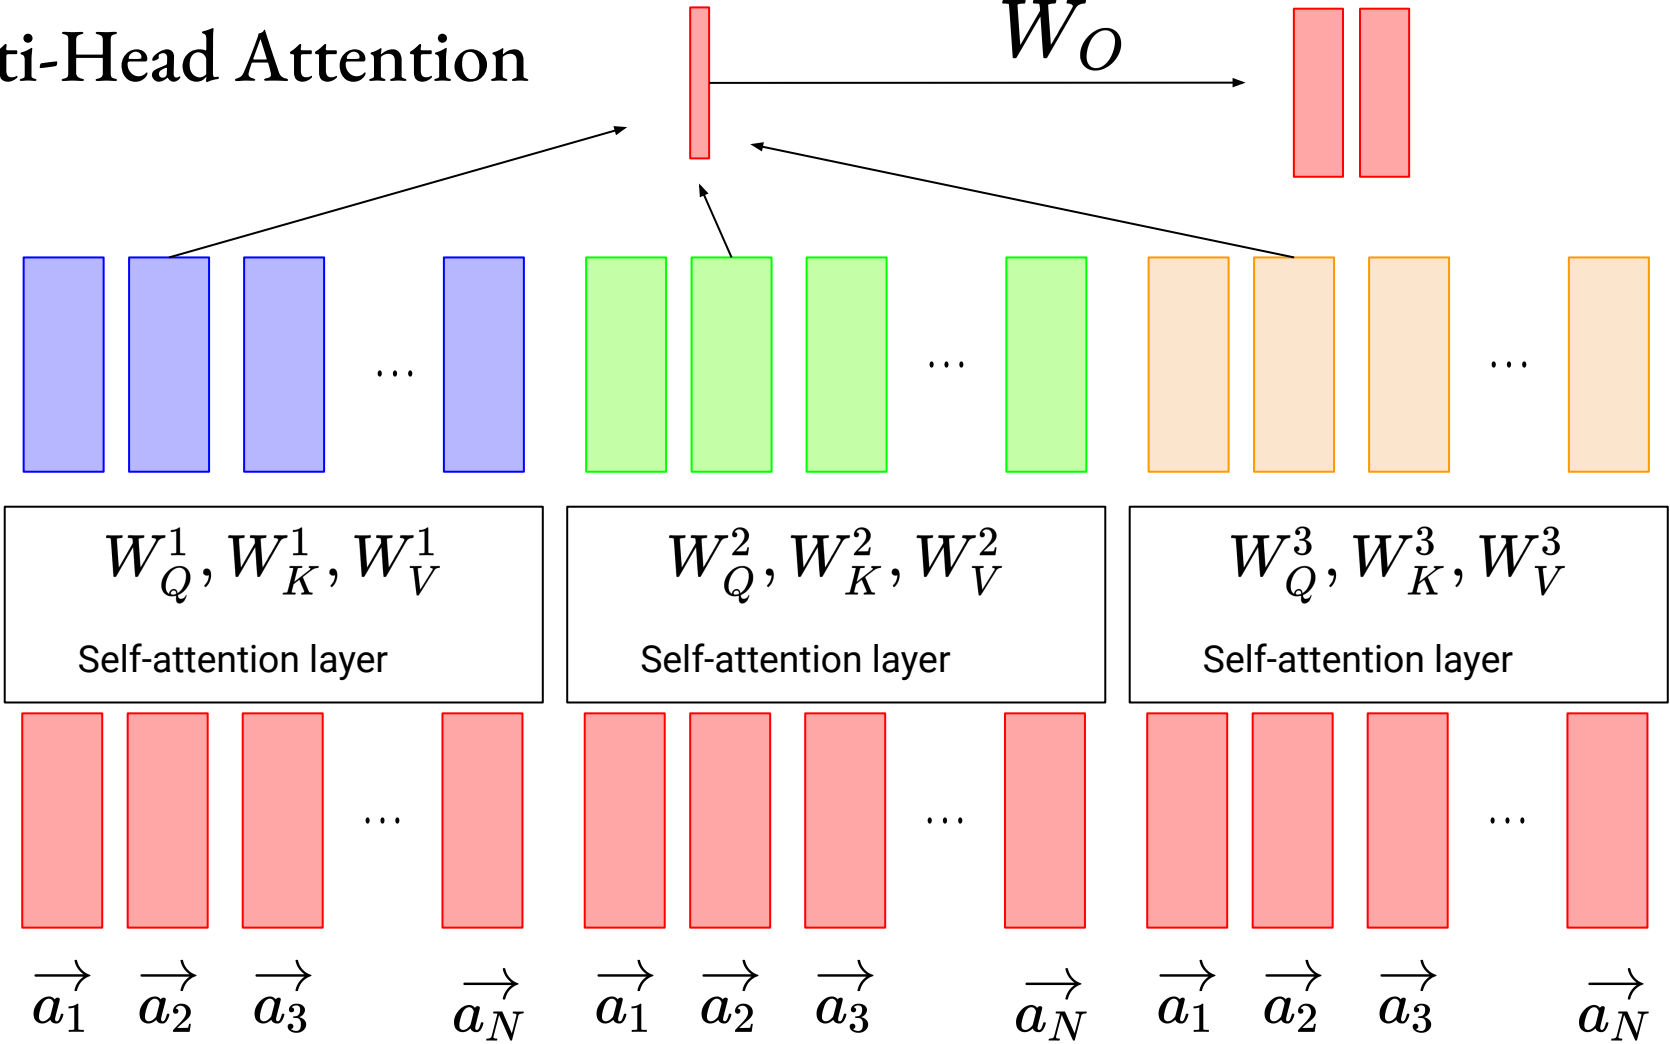
# Multi-Head Attention



# Multi-Head Attention

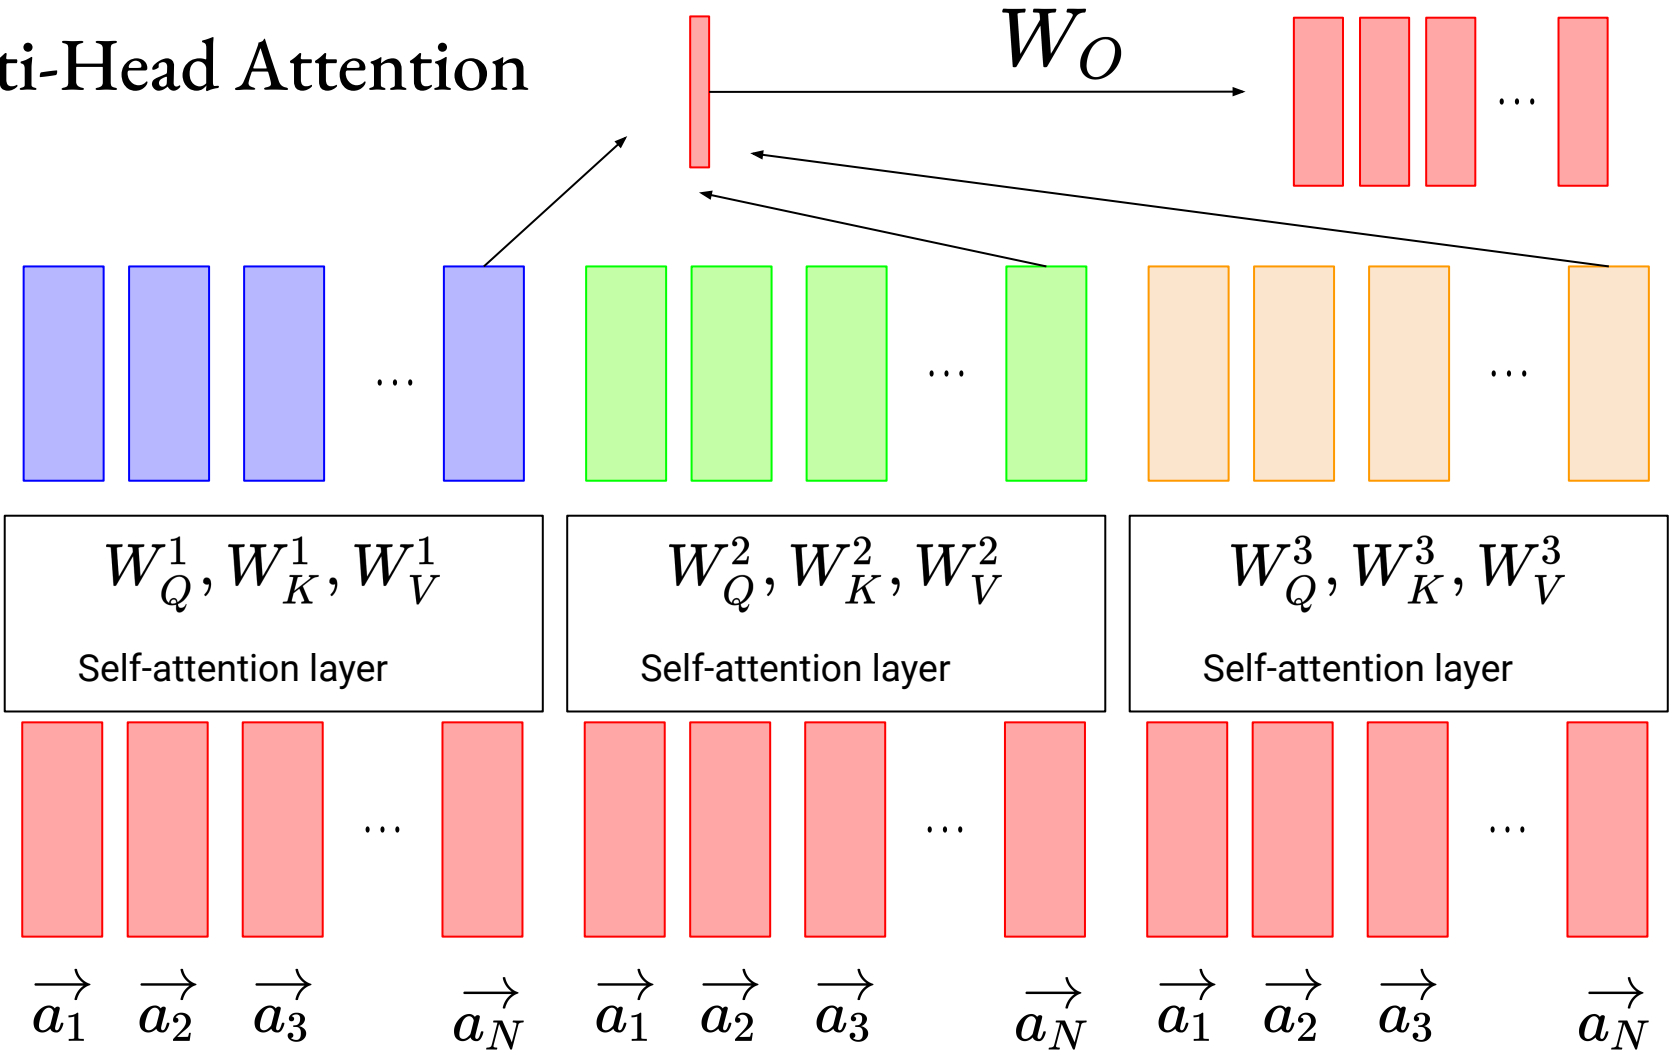


# Multi-Head Attention






# Multi-Head Attention



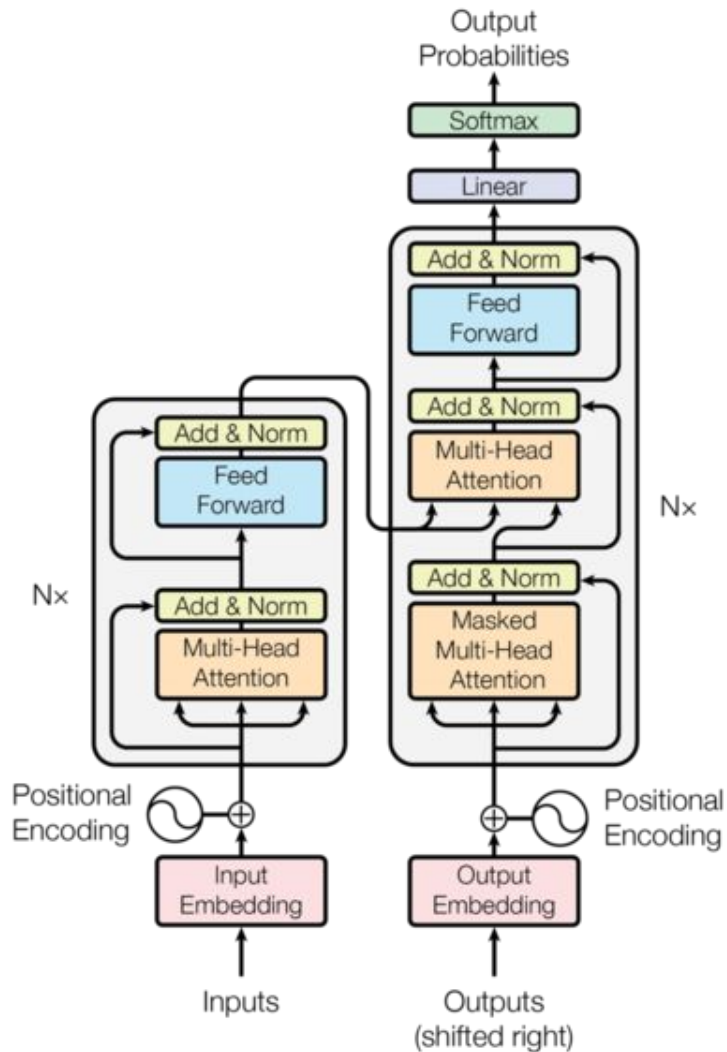
# Transformer

- You can now understand everything in this famous diagram from the “Attention is All You Need” paper.



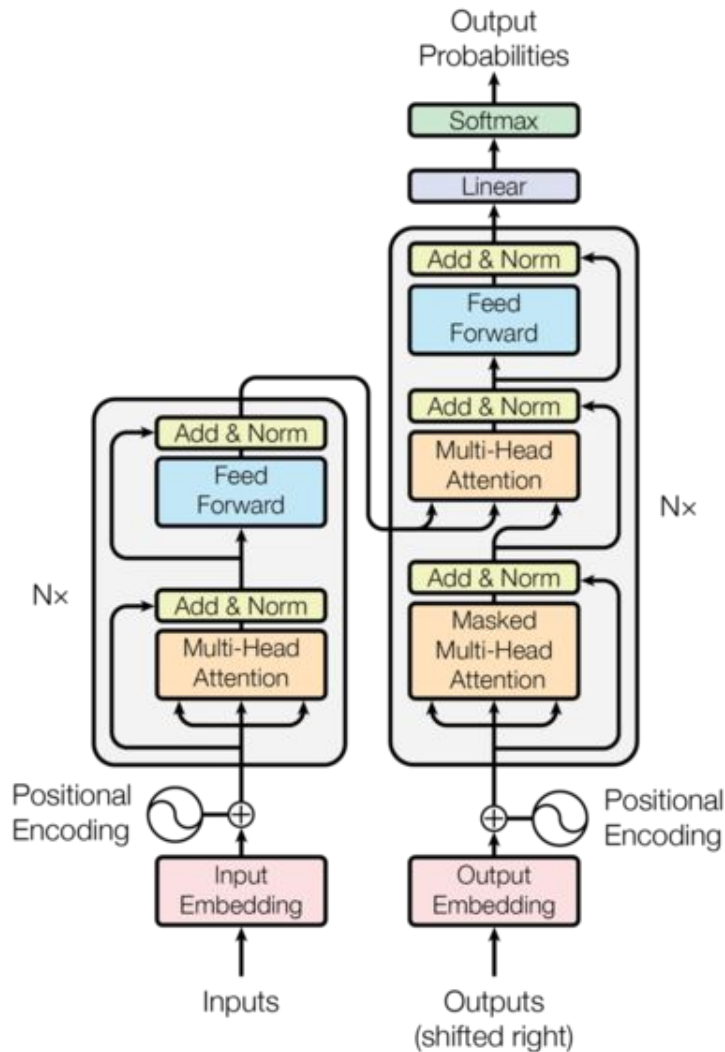
# Transformer

- You can now understand everything in this famous diagram from the “Attention is All You Need” paper.
  1. Add the initial sequence embedding to the output of the MHA layer.
  2. Pass it through a Layer Normalization.




# Transformer

- You can now understand everything in this famous diagram from the “Attention is All You Need” paper.
- “Add & Norm” is like a residual connection
- That output gets put into a feed forward layer (and then another residual connection)

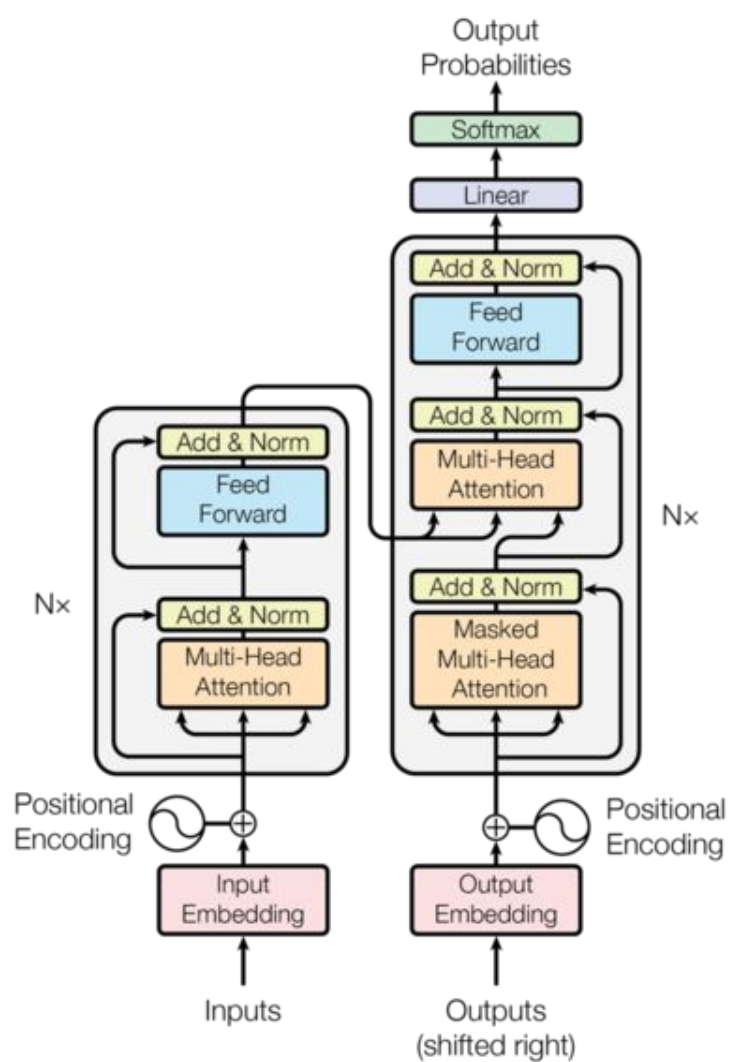



# Transformer

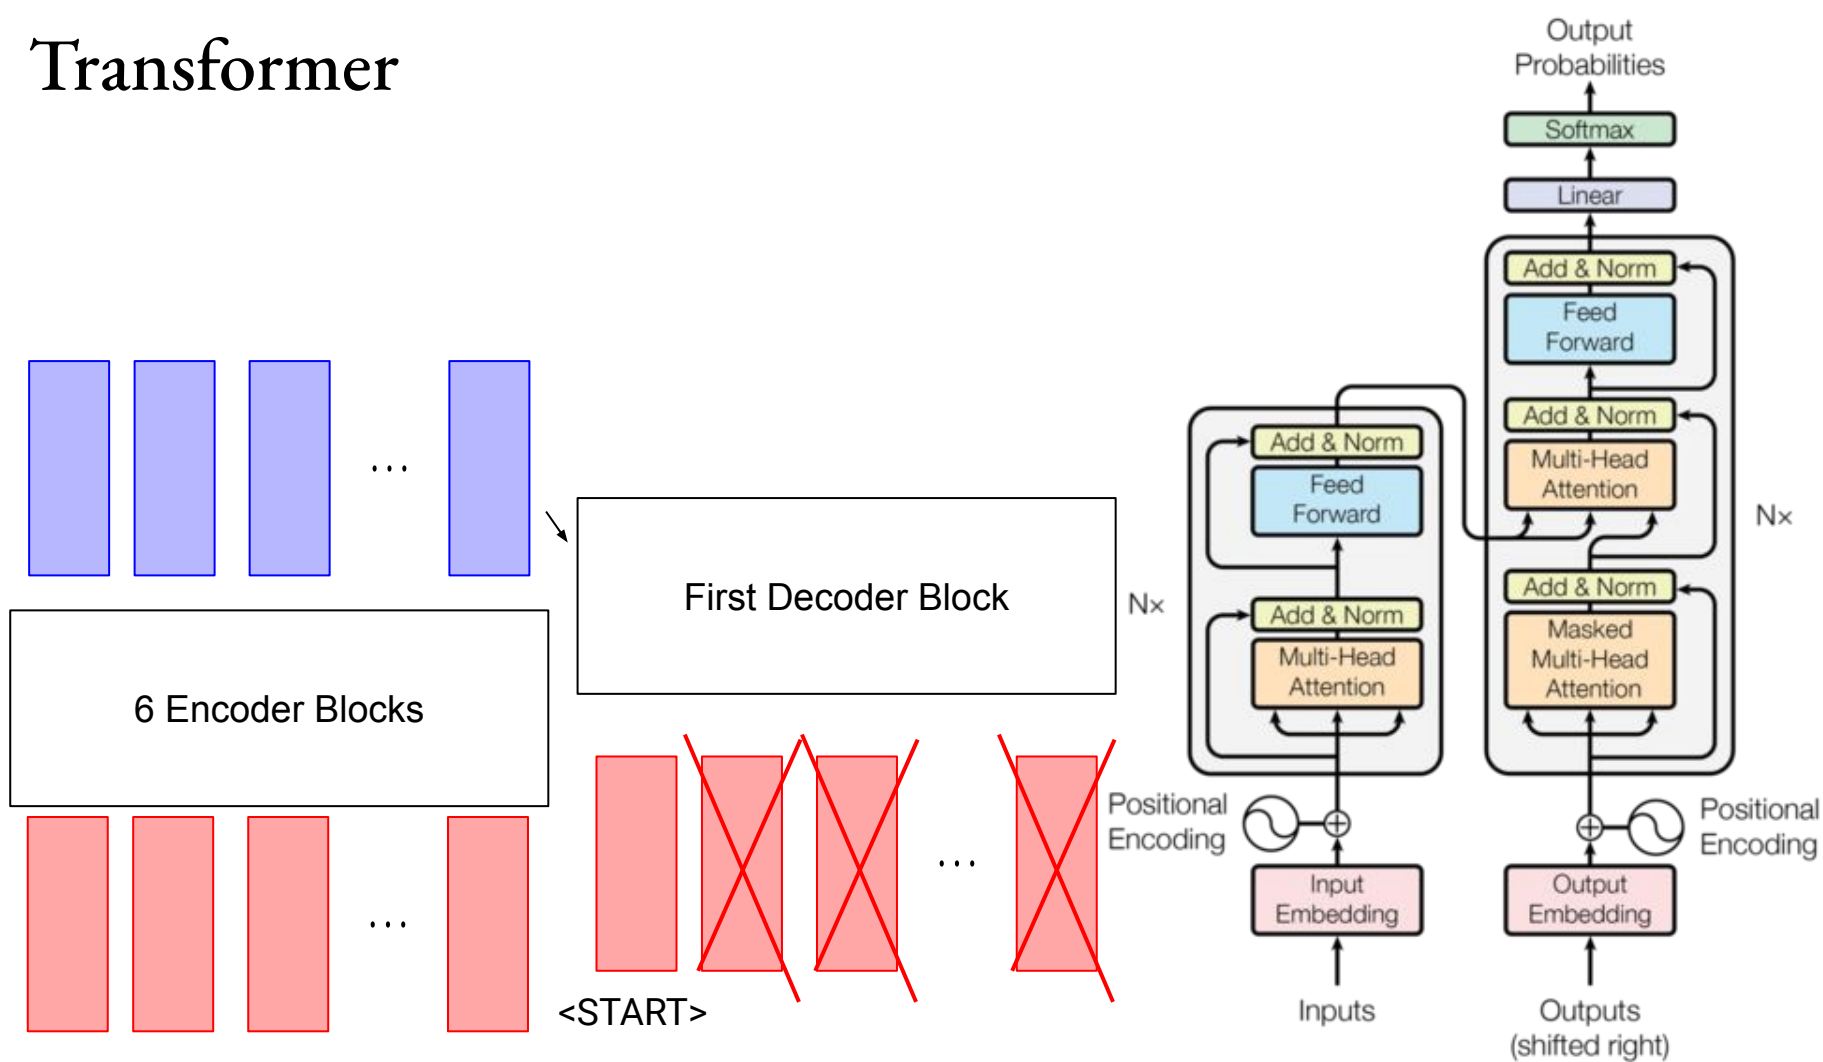
- You can now understand everything in this famous diagram from the “Attention is All You Need” paper.
- “Add & Norm” is like a residual connection
- That output gets put into a feed forward layer (and then another residual connection)
- Original model has 6 Transformer encoder layers stacked on one another.



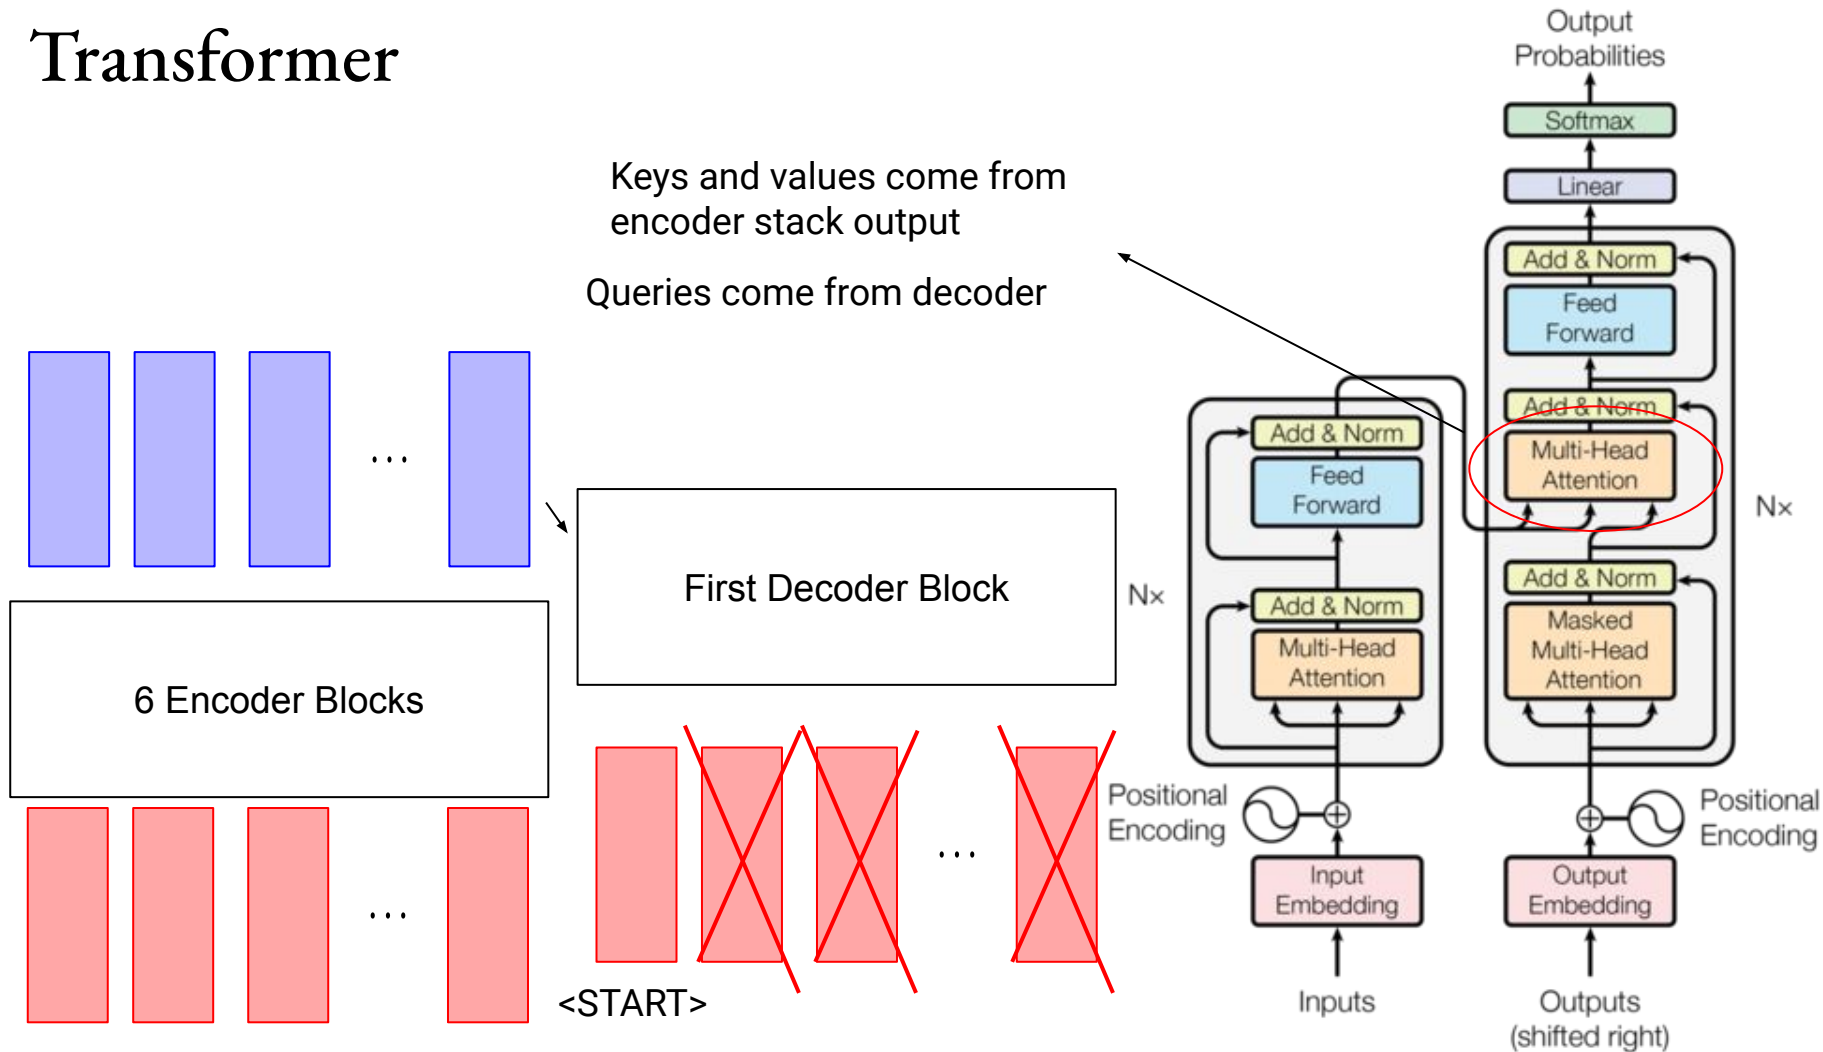
# Transformer



# Transformer

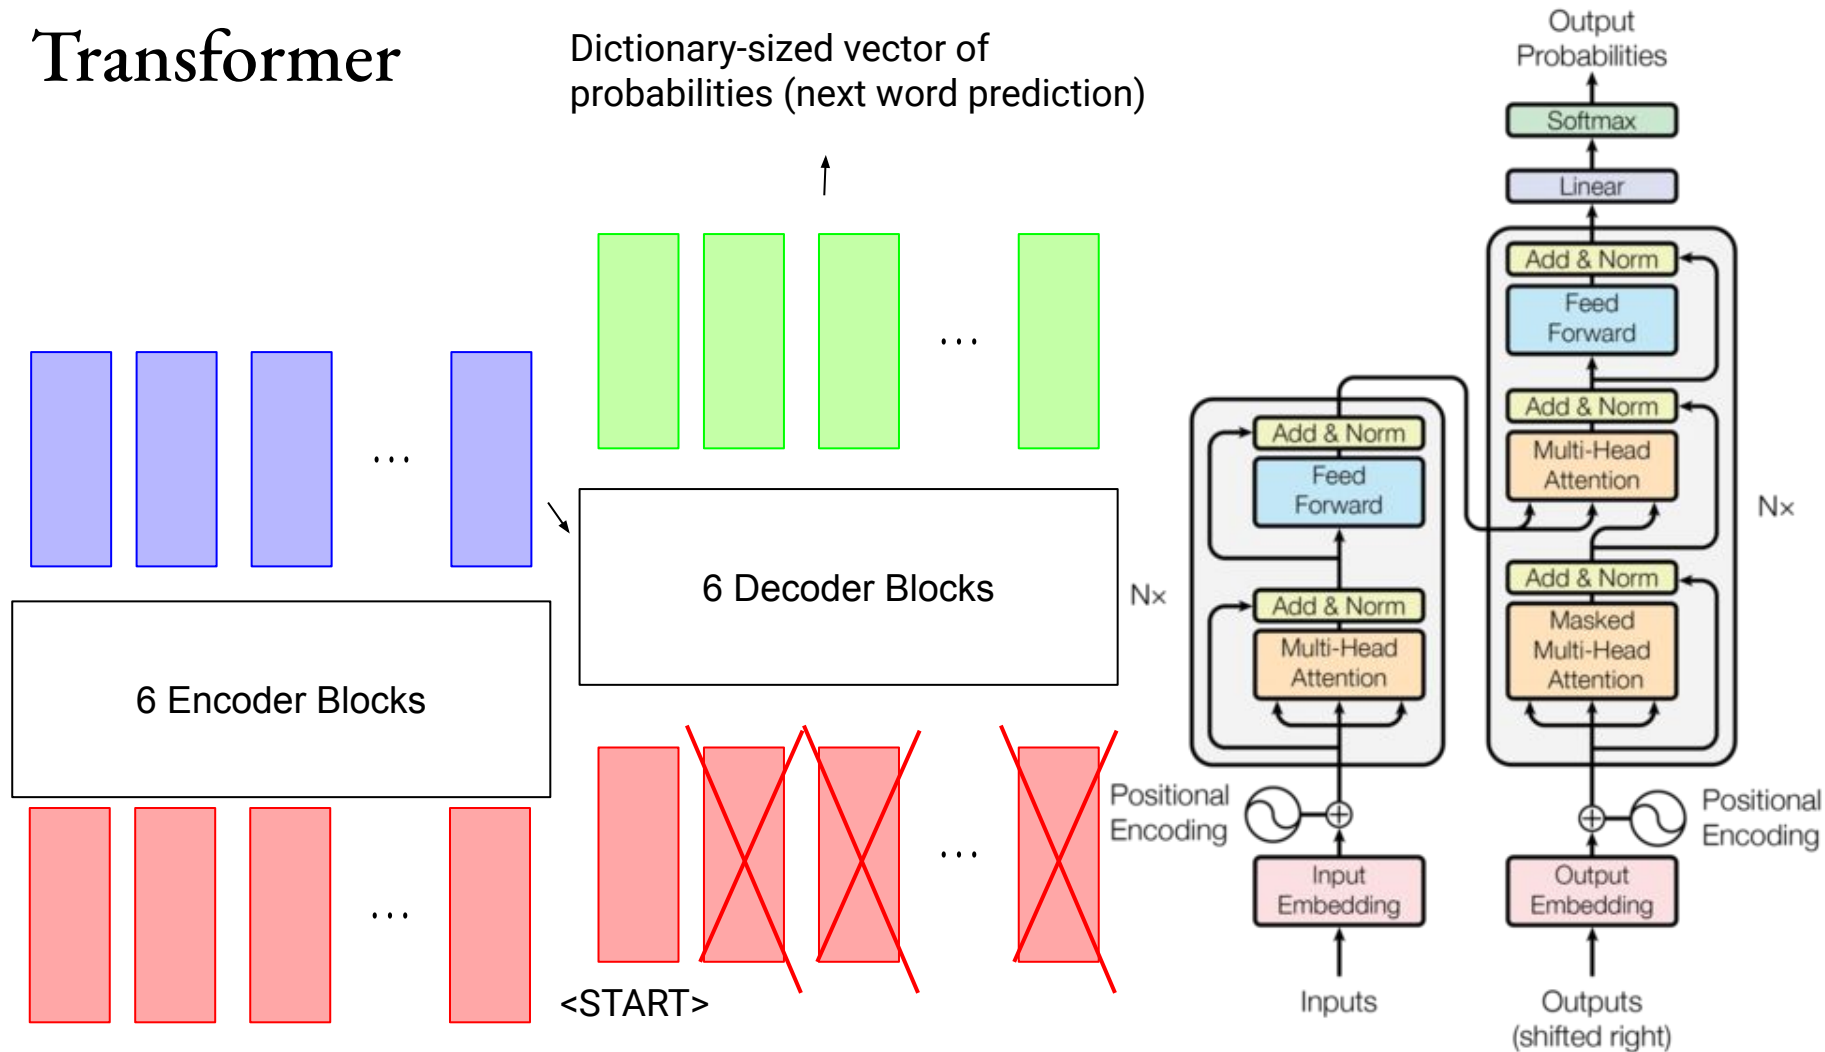


# Transformer

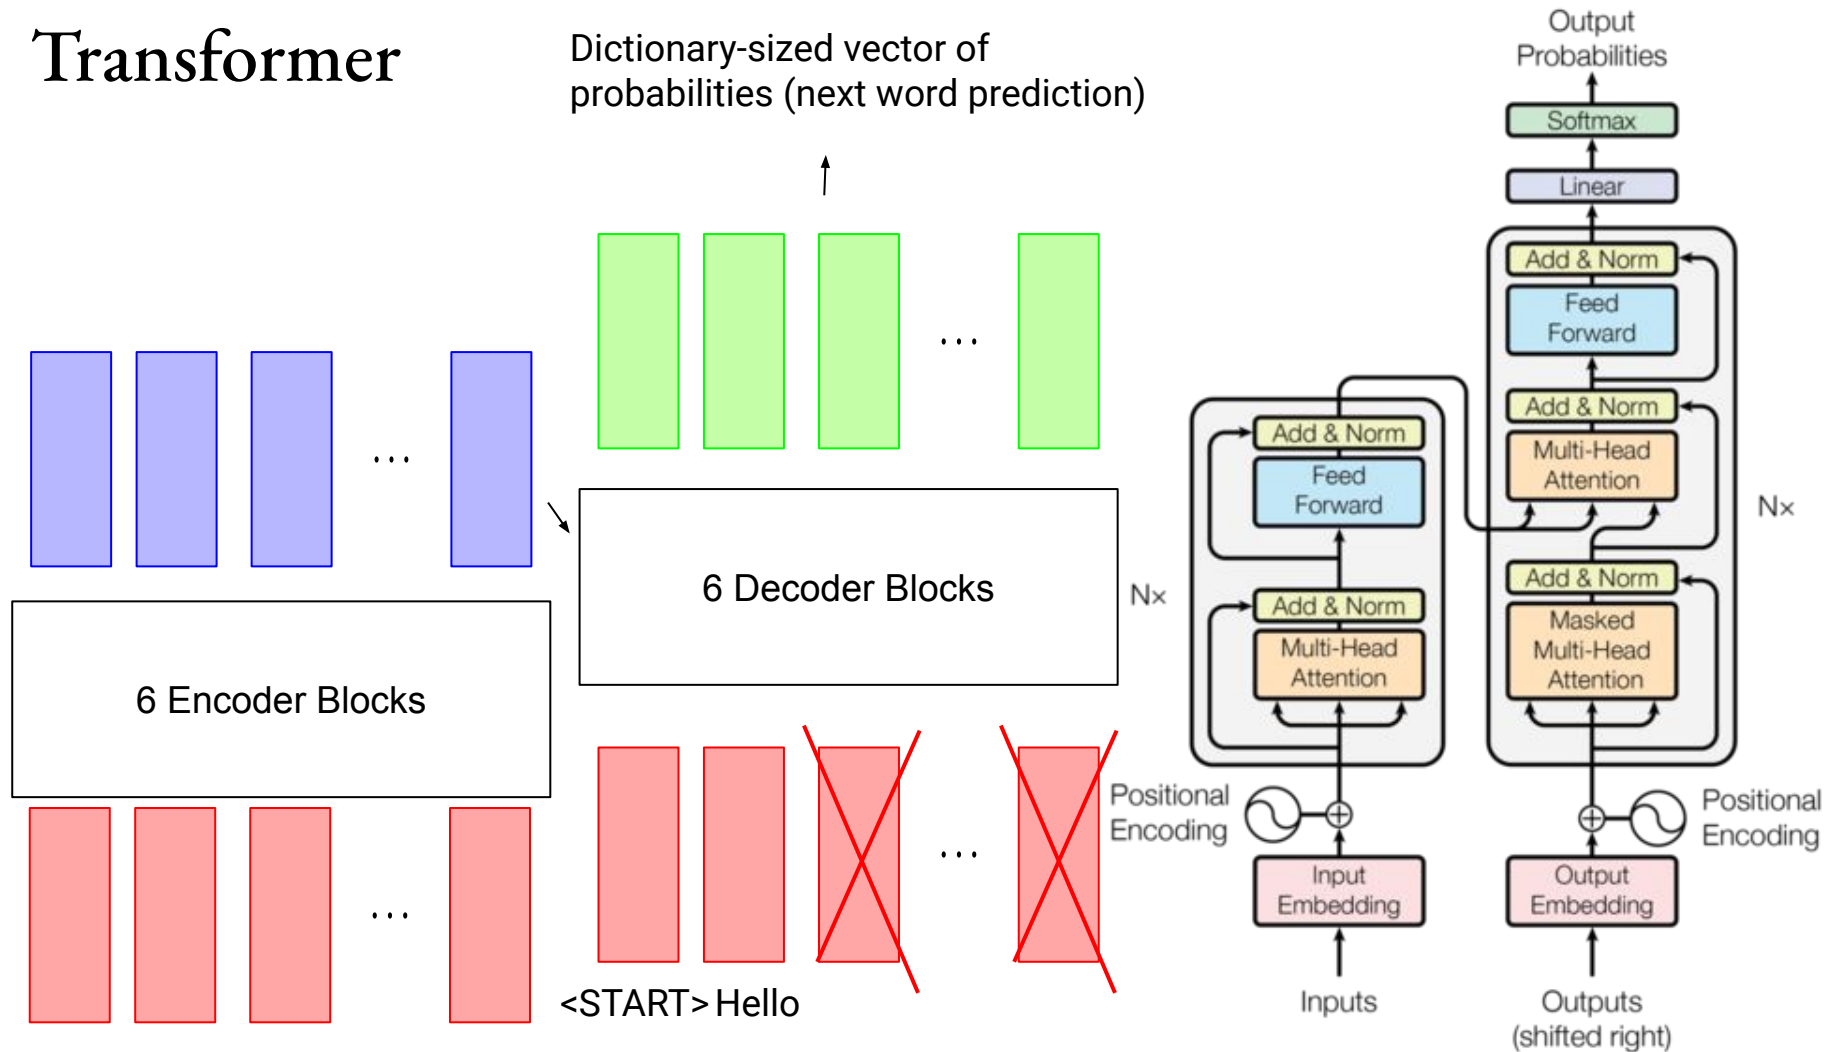




# Transformer




# Transformer




# Self-Supervised Learning

- BERT: Bidirectional Encoder Representations from Transformers
  - Encoder-only model
  - 12 Encoders



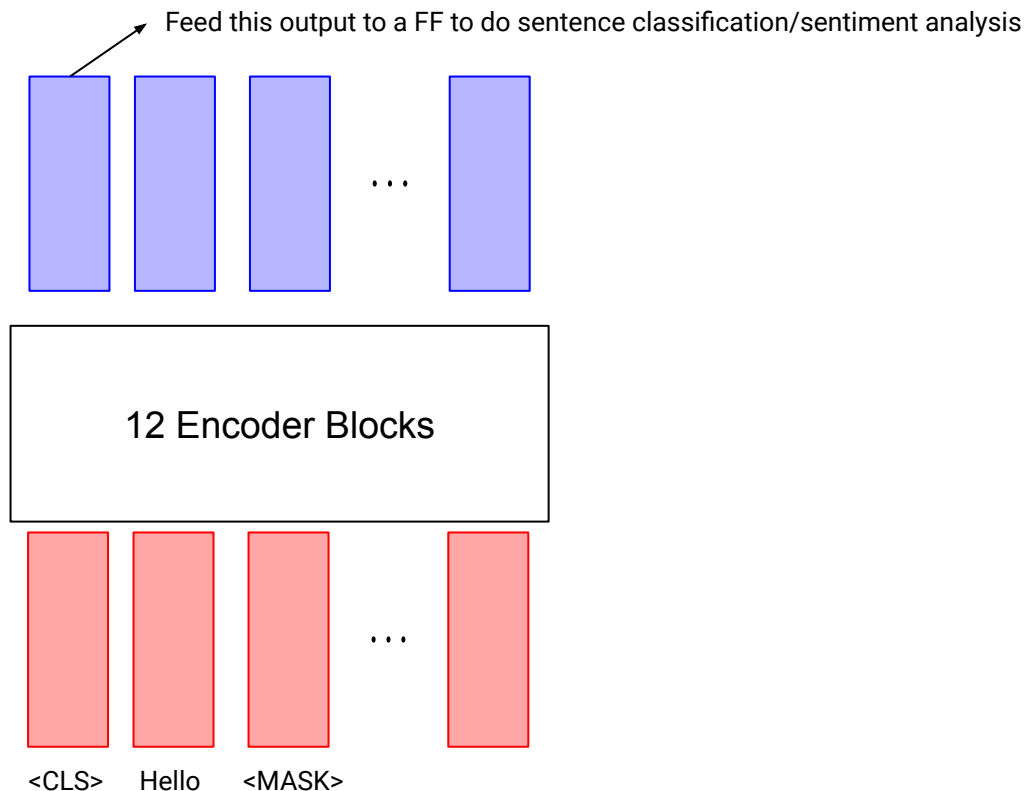
# Self-Supervised Learning

- BERT: Bidirectional Encoder Representations from Transformers
  - Encoder-only model
  - 12 Encoders
- Self-supervised Task #1
  1. Take a mountain of text.
  2. Mask ~15% of the words
  3. Try to predict the masked words.



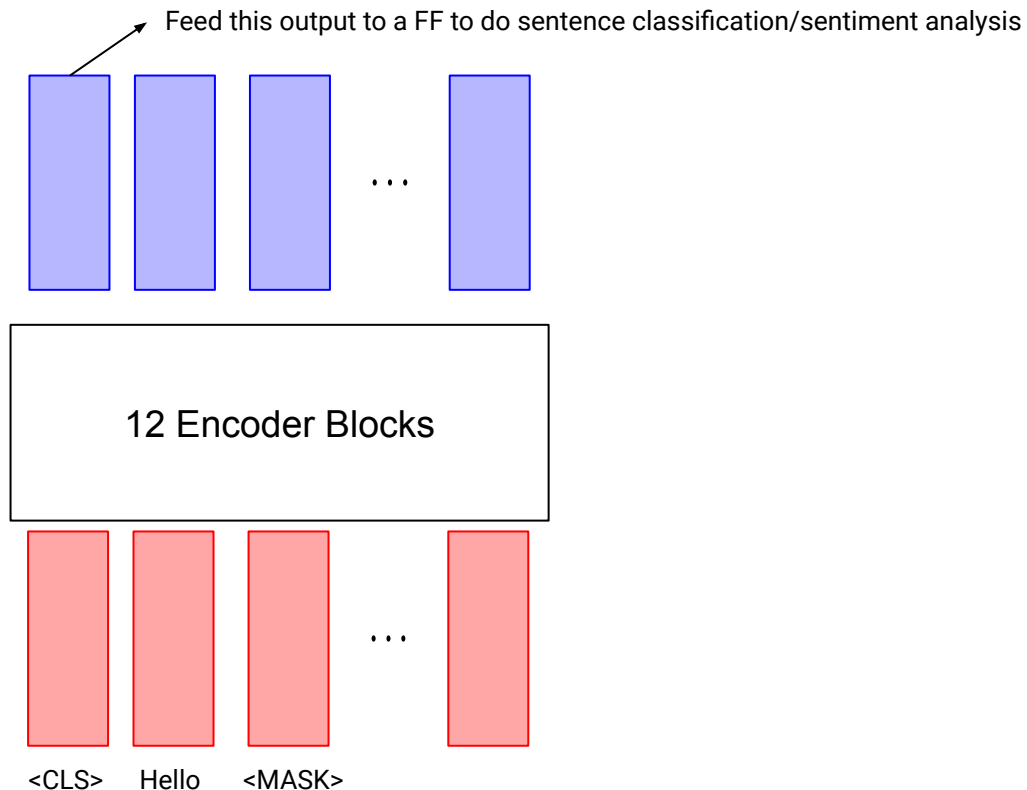
# Self-Supervised Learning

- BERT: Bidirectional Encoder Representations from Transformers
  - Encoder-only model
  - 12 Encoders
- Self-supervised Task #1
  1. Take a mountain of text.
  2. Mask ~15% of the words
  3. Try to predict the masked words.
- <CLS> is a special token



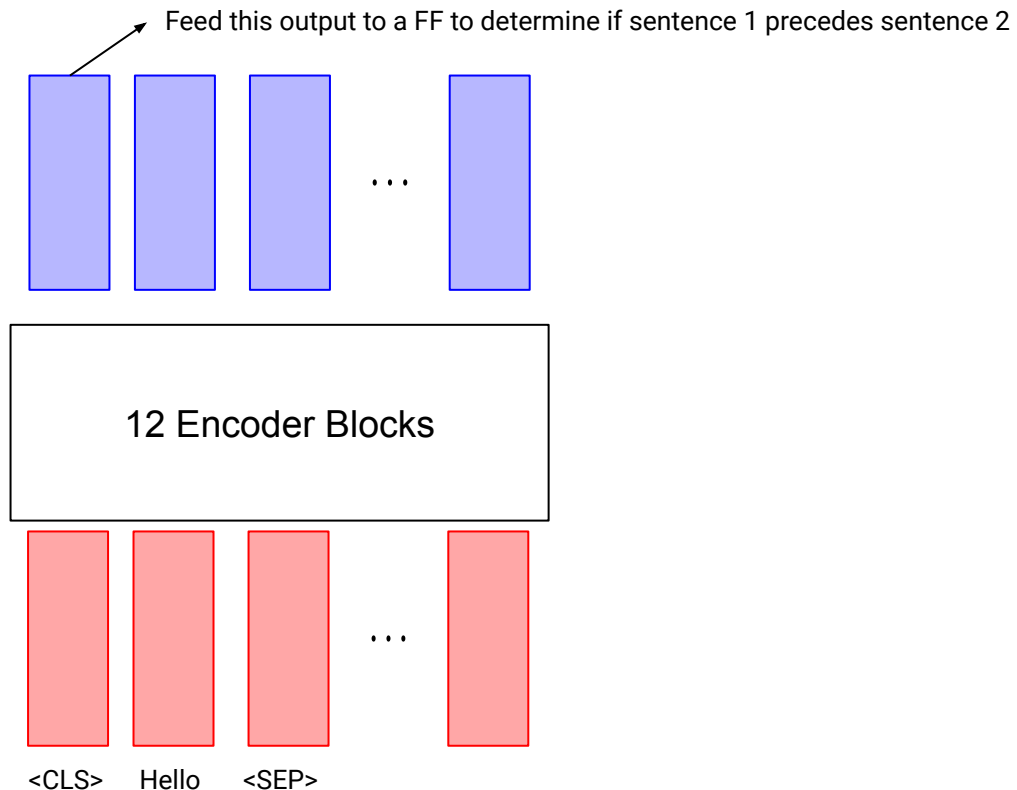
# Self-Supervised Learning

- BERT: Bidirectional Encoder Representations from Transformers
  - Encoder-only model
  - 12 Encoders
- Self-supervised Task #1
  1. Take a mountain of text.
  2. Mask ~15% of the words
  3. Try to predict the masked words.
- <CLS> is a special token
- <SEP> is a token that separates sentences (Q/A tasks)




# Self-Supervised Learning

- BERT: Bidirectional Encoder Representations from Transformers
  - Encoder-only model
  - 12 Encoders
- Self-supervised Task #2
  1. Take a mountain of text.
  2. Take Two Sentences
  3. Predict whether the first sentence follows the second
- <CLS> is a special token
- <SEP> is a token that separates sentences (Q/A tasks)



# Self-Supervised Learning

- BERT: Bidirectional Encoder Representations from Transformers
  - Encoder-only model
  - 12 Encoders
- Fine-tune the pre-trained BERT for your task of choice!
- Can use frozen embeddings from any layer of BERT! (feature extraction)





# Nice Blog Posts

- [Illustrated Transformer](#)
- [Illustrated BERT](#)